

Pipelined MIPS Processor

Dr. Igor Ivkovic

iivkovic@uwaterloo.ca

[with material from “Computer Organization and Design” by Patterson and Hennessy, and “Digital Design and Computer Architecture” by Harris and Harris, both published by Morgan Kaufmann]

Objectives

- Introduction to MIPS Pipeline
- Different Hazard Types
- Handling Data Hazards
- Handling Control Hazards
- Handling Exceptions
- Advanced MIPS Pipeline Design

Introduction to MIPS Pipeline /1

■ **Pipelining Introduced:**

- The goal is to improve the throughput of digital circuitry
- We design a MIPS pipeline by dividing the single-cycle processor into five pipeline stages
- In each stage, one instruction can run, so five instructions can run simultaneously
- **Pipelining can ideally improve the throughput by five times**
- Pipelining does introduce overhead, but unlike the multi-cycle processor, the performance advantage is significant
- **As the result, most modern high-performance digital architectures are based on pipelining**

Introduction to MIPS Pipeline /2

■ **Pipelining Elaborated:**

- Similar as in the multi-cycle processor, we focus on the operations that create the largest time delay
- These operations include reading from and writing to memory, register file access, and ALU operations
- Five stages, one step per stage

■ **The stages are similar to the stages that multi-cycle processor goes through when executing `lw` :**

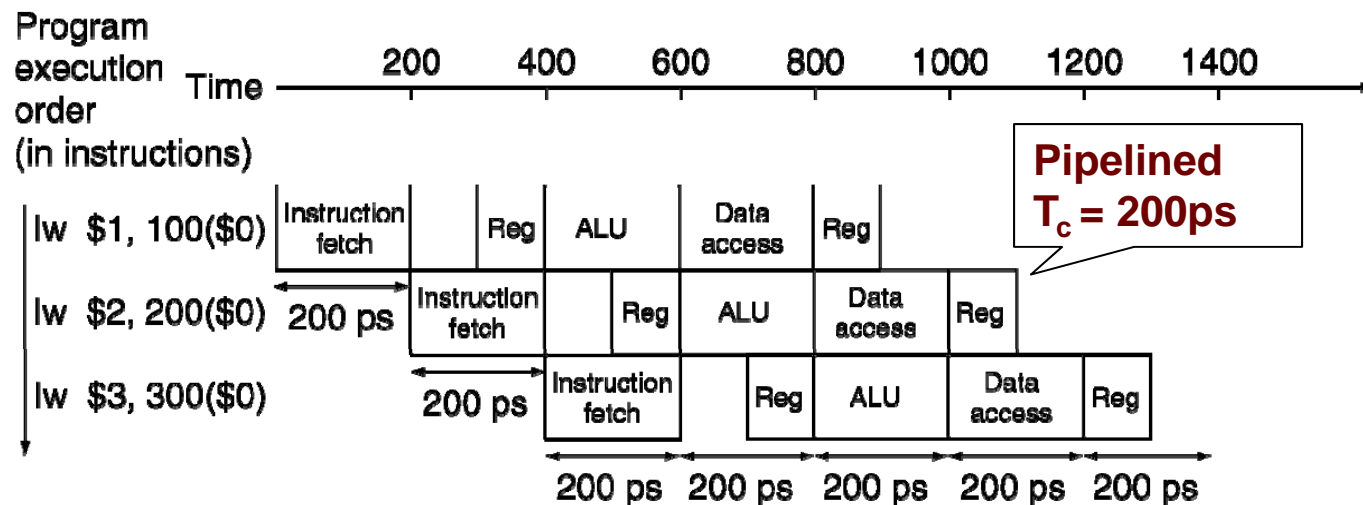
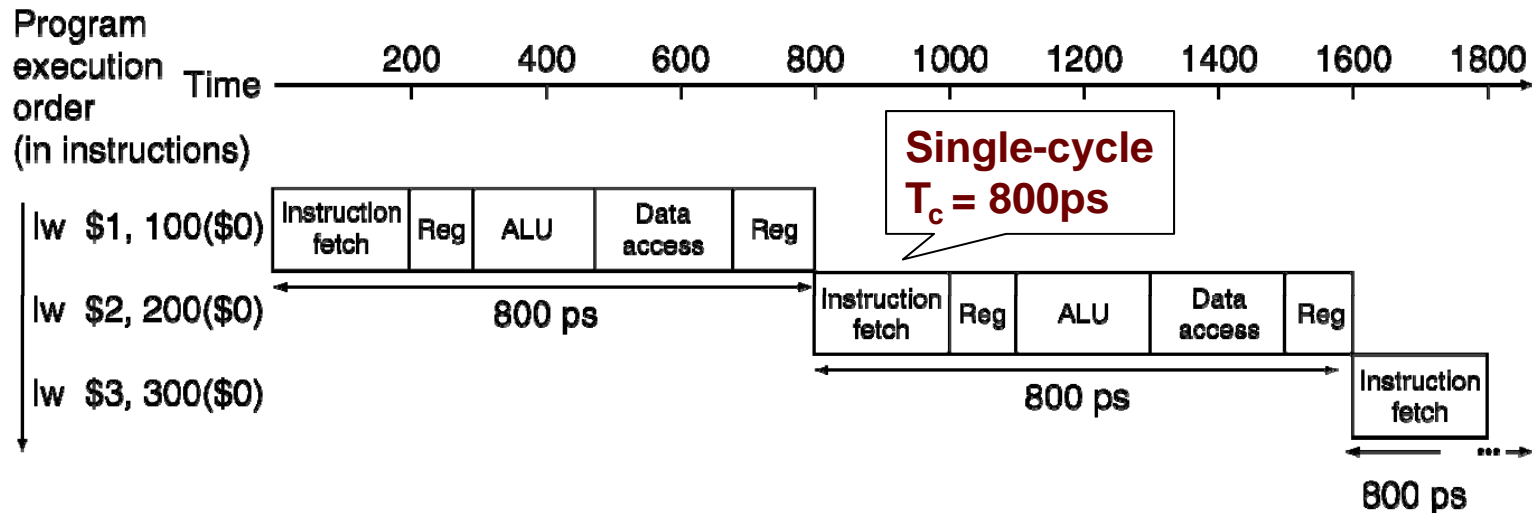
- **Fetch (IF):** Instruction fetch from memory
- **Decode (ID):** Instruction decode and register read
- **Execute (EX):** Execute operation or calculate address
- **Memory (MEM):** Access memory operand
- **Writeback (WB):** Write result back to register

Introduction to MIPS Pipeline /3

- **Five stages, one step per stage:**
 - **Fetch (IF):** Instruction fetch from memory
 - **Decode (ID):** Instruction decode and register read
 - **Execute (EX):** Execute operation or calculate address
 - **Memory (MEM):** Access memory operand
 - **Writeback (WB):** Write result back to register
- **Add pipeline registers between stages**
- **Pipeline performance:**
 - Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages
 - Let us compare pipelined datapath with single-cycle datapath

Introduction to MIPS Pipeline /4

■ Pipeline Performance:



Introduction to MIPS Pipeline /5

■ Pipeline Speedup:

- If all stages are balanced and take the same time, then

{time between instructions}_{pipelined} =

{time between instructions}_{non-pipelined} / {number of stages}

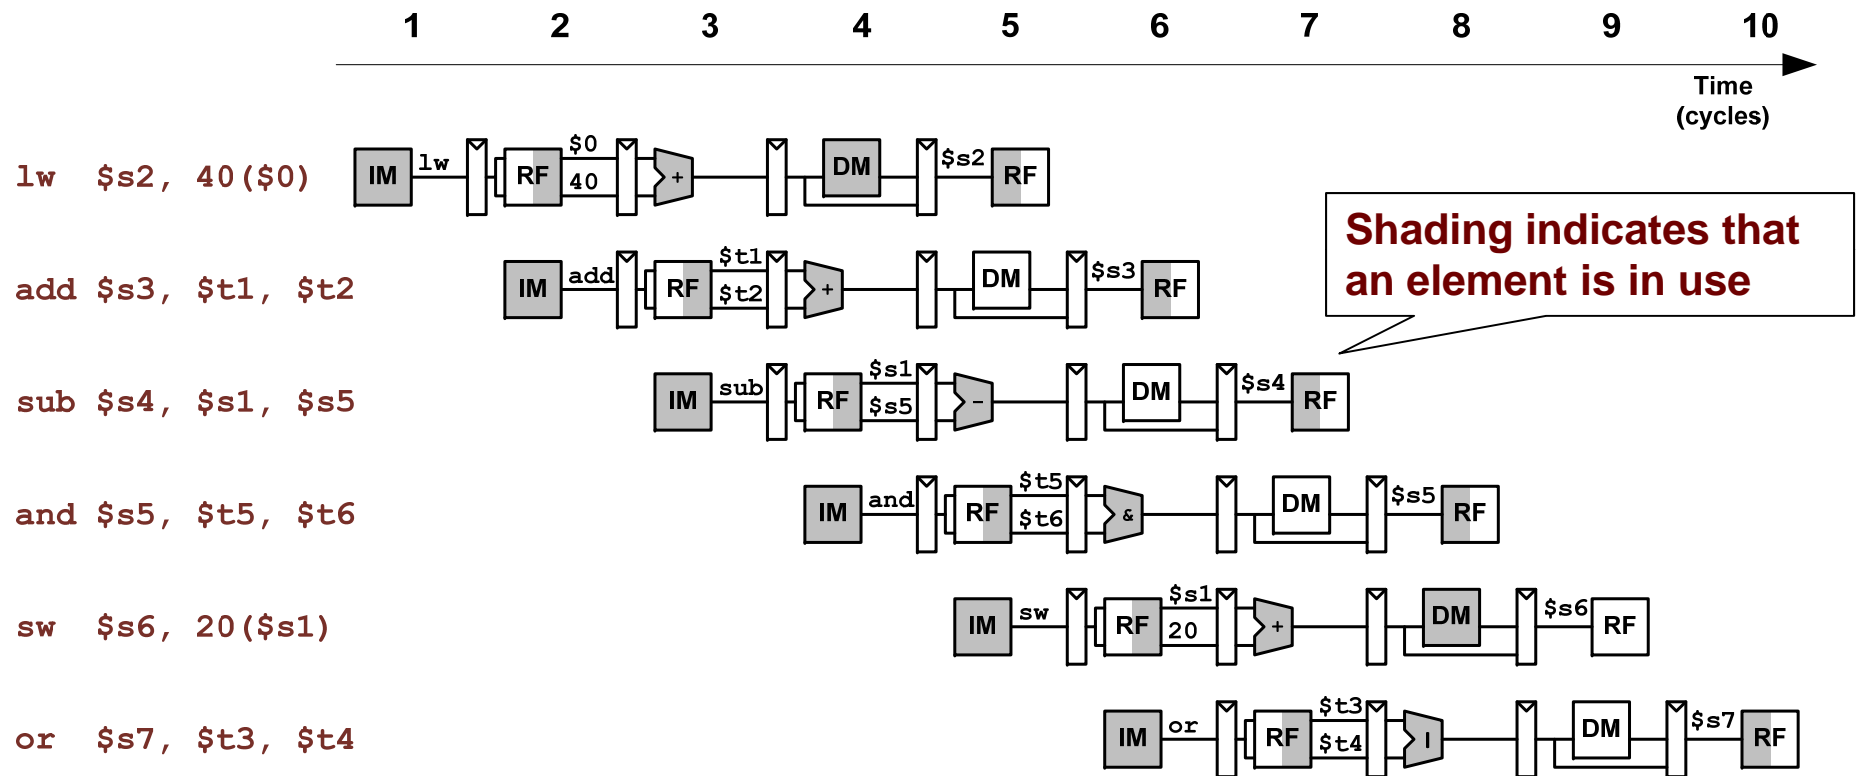
- The speedup is due to the increased throughput
- **If not balanced, the pipeline speedup is less**
- Latency per each instruction does not decrease

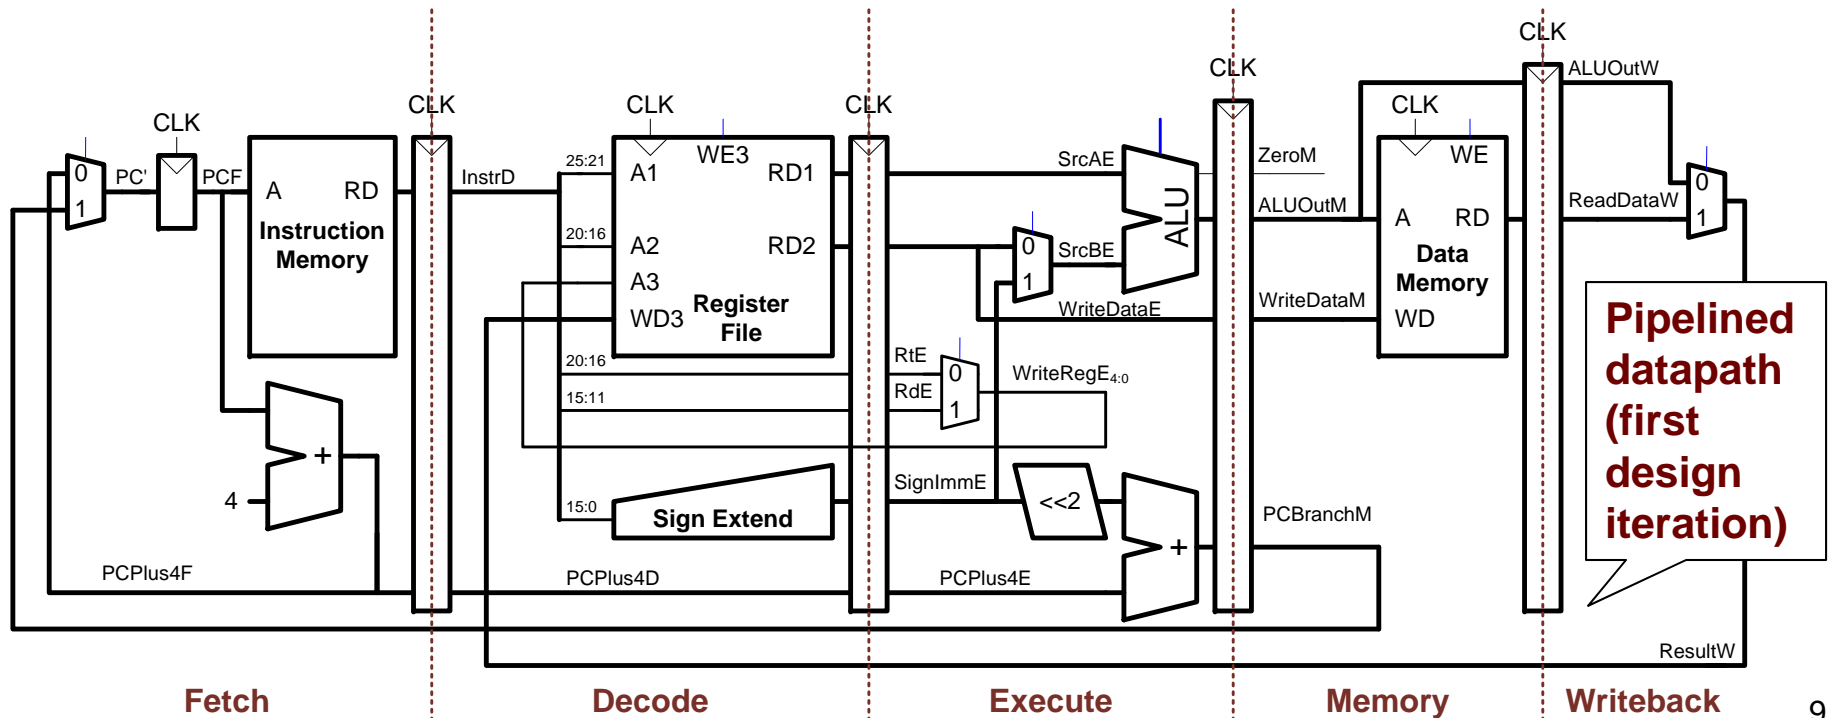
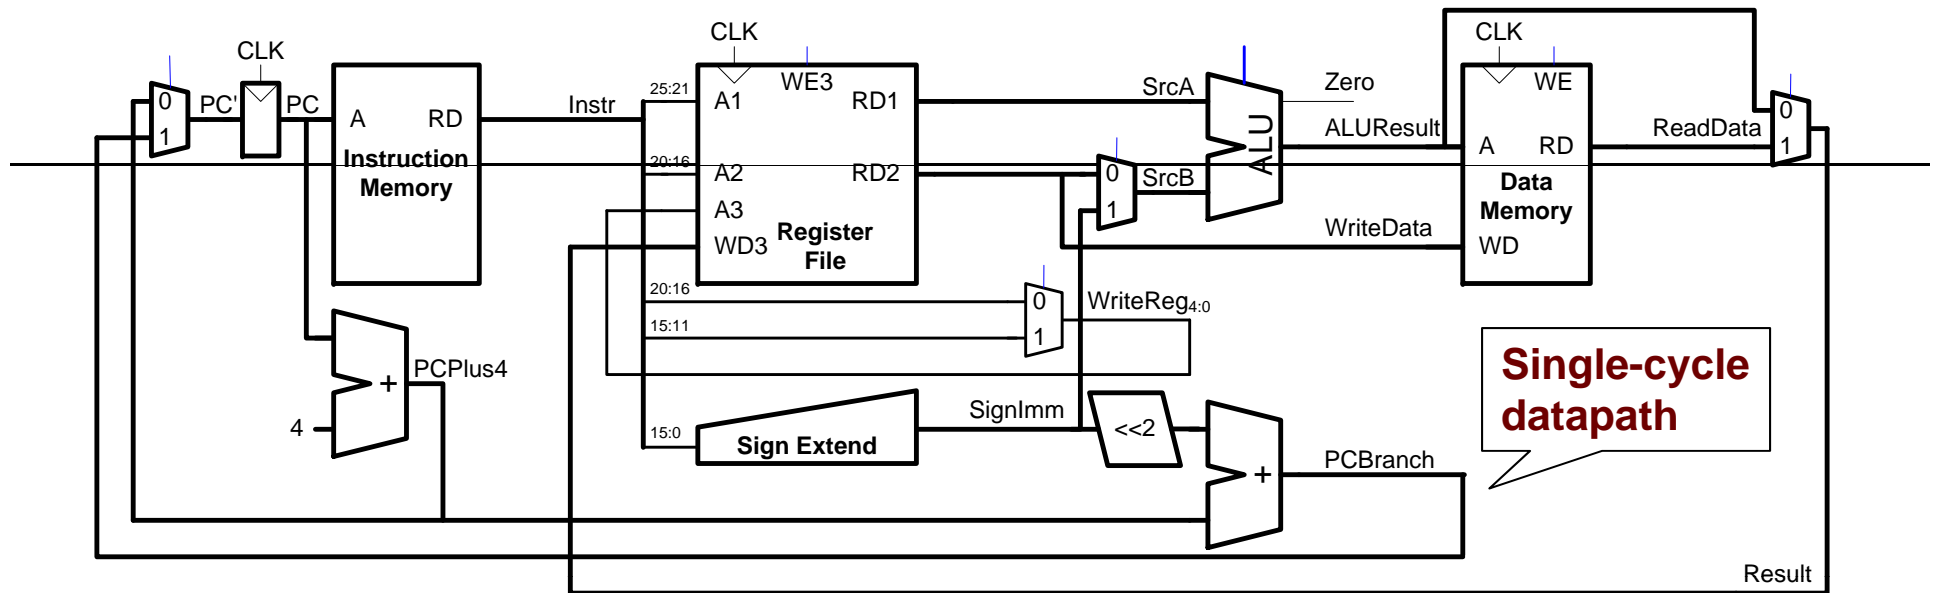
■ MIPS ISA designed for pipelining:

- All instructions are 32-bits, so they are easier to fetch and decode in one cycle
- MIPS allows decoding and reading of registers in one step
- Memory operands are aligned, so memory access takes only one cycle

Introduction to MIPS Pipeline /6

■ Pipelining Elaborated:

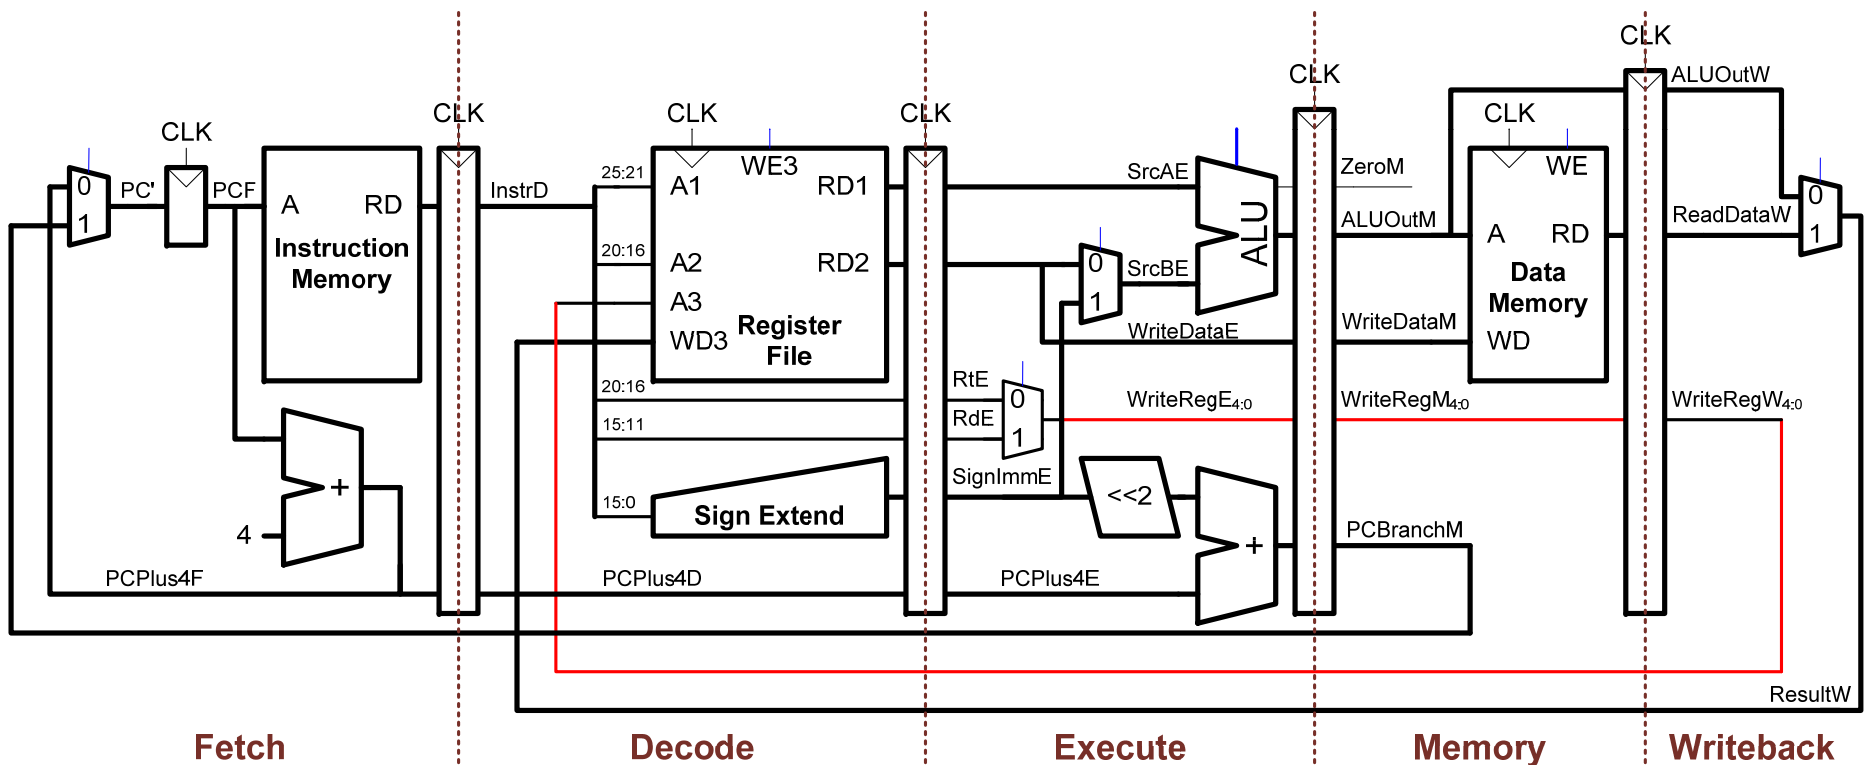




Introduction to MIPS Pipeline /8

■ Corrected Pipelined Datapath:

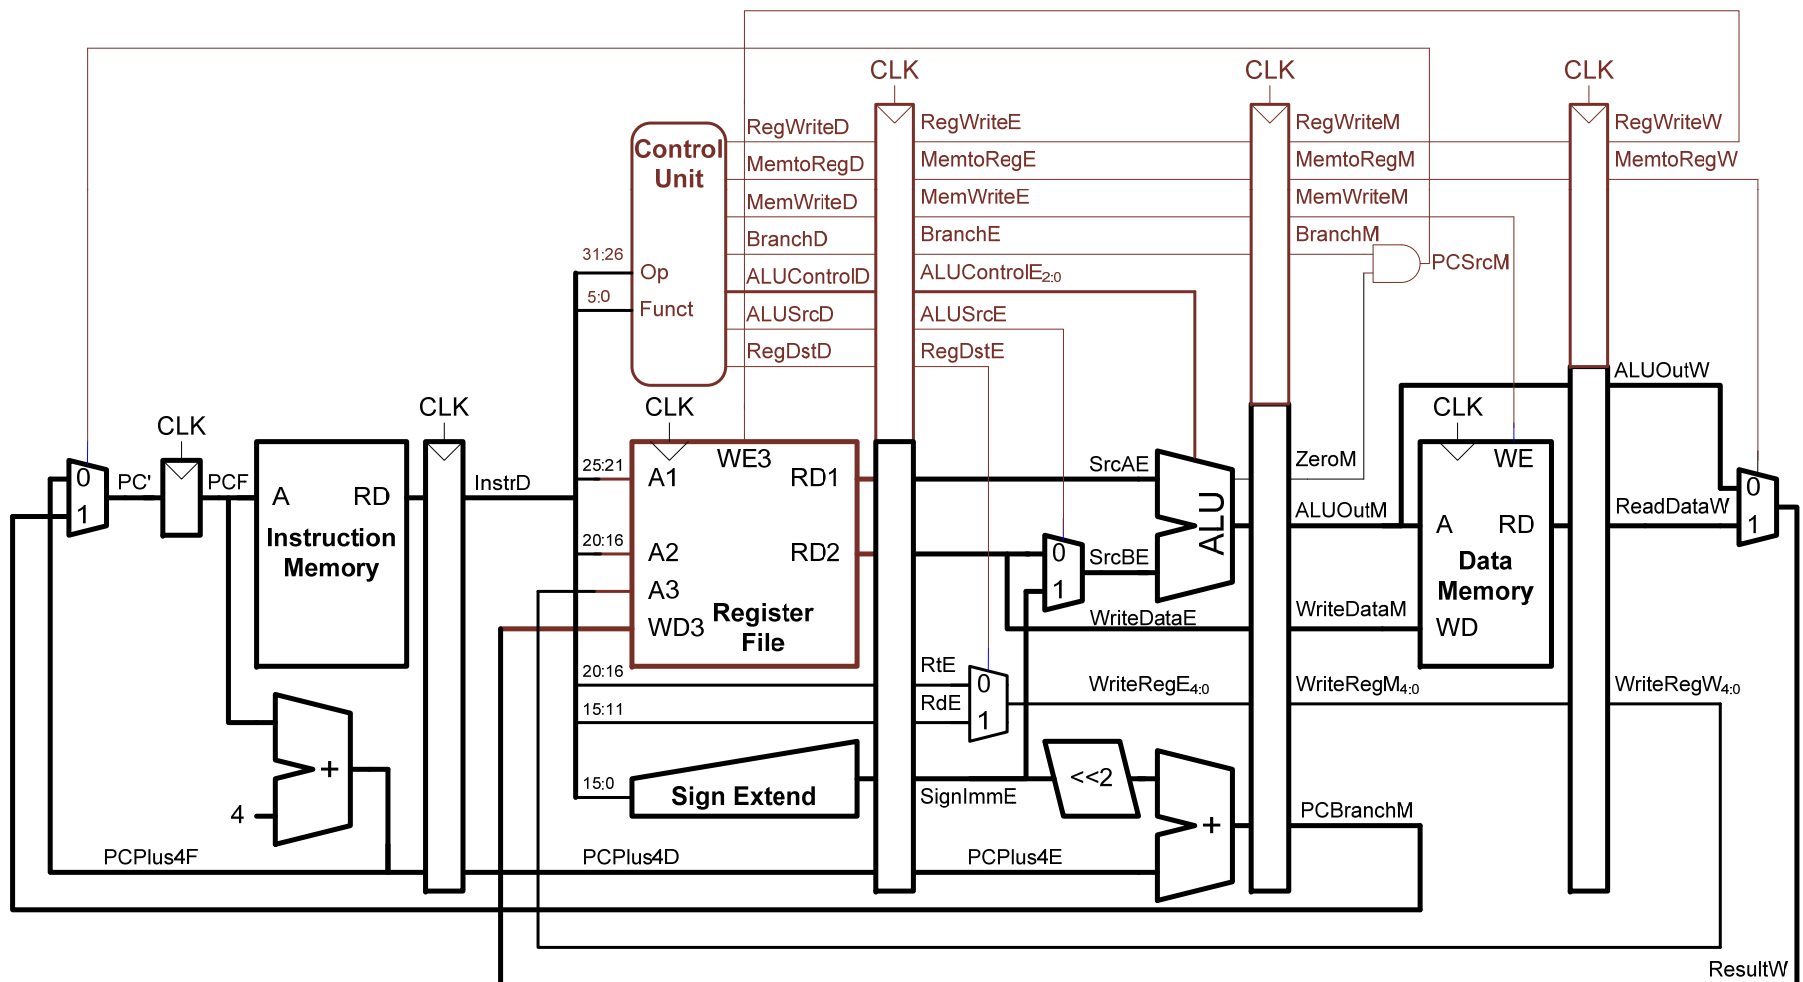
- WriteReg must be carried through and arrive at the same time as the Result signal; otherwise, lw on cycle 5 writes to \$s4



Introduction to MIPS Pipeline /9

■ Pipelined Datapath Control Unit:

- Same control unit as single-cycle processor
- Control delayed to proper pipeline stage



Introduction to MIPS Pipeline /10

■ Hazard:

- A problem in execution where the subsequent instruction cannot execute
- Occurs when an instruction depends on result from instruction that hasn't completed

■ Types of Hazards:

- **Structure Hazard:** A required architecture resource is busy
 - For instance, instruction fetch would have to stall for that cycle
 - **Addressed by separating Instruction and Data Memory into different units**
- **Data Hazard:** Need to wait for previous instruction to complete its data read/write
- **Control Hazard:** Next instruction not decided yet
 - Typically caused by branching
 - Predict outcome of branch and only stall if prediction is wrong

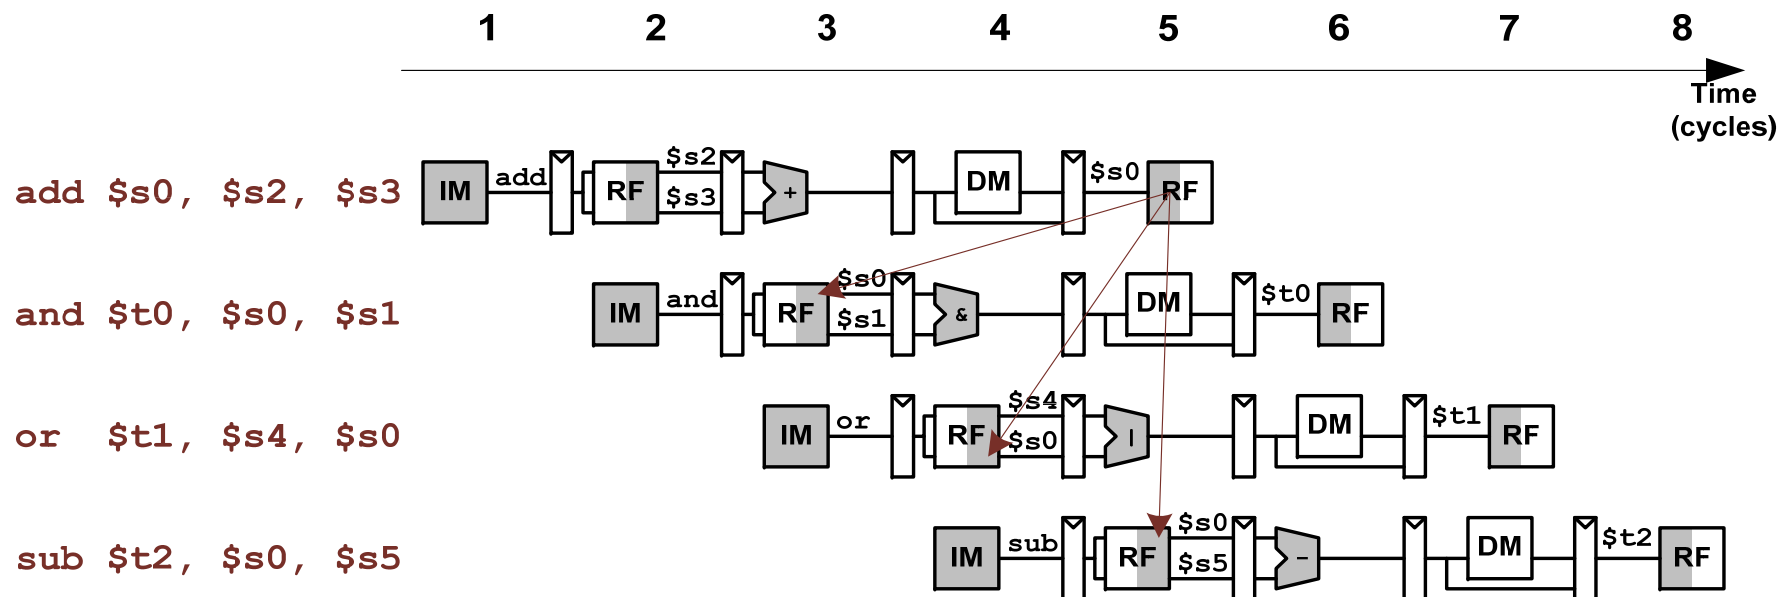
Handling Data Hazards /1

■ Data Hazards:

- Register value not yet written back to register file

■ Read After Write (RAW) Hazard:

- Read occurs after the needed Write



and and or obtain incorrect values,
but sub obtains the correct value

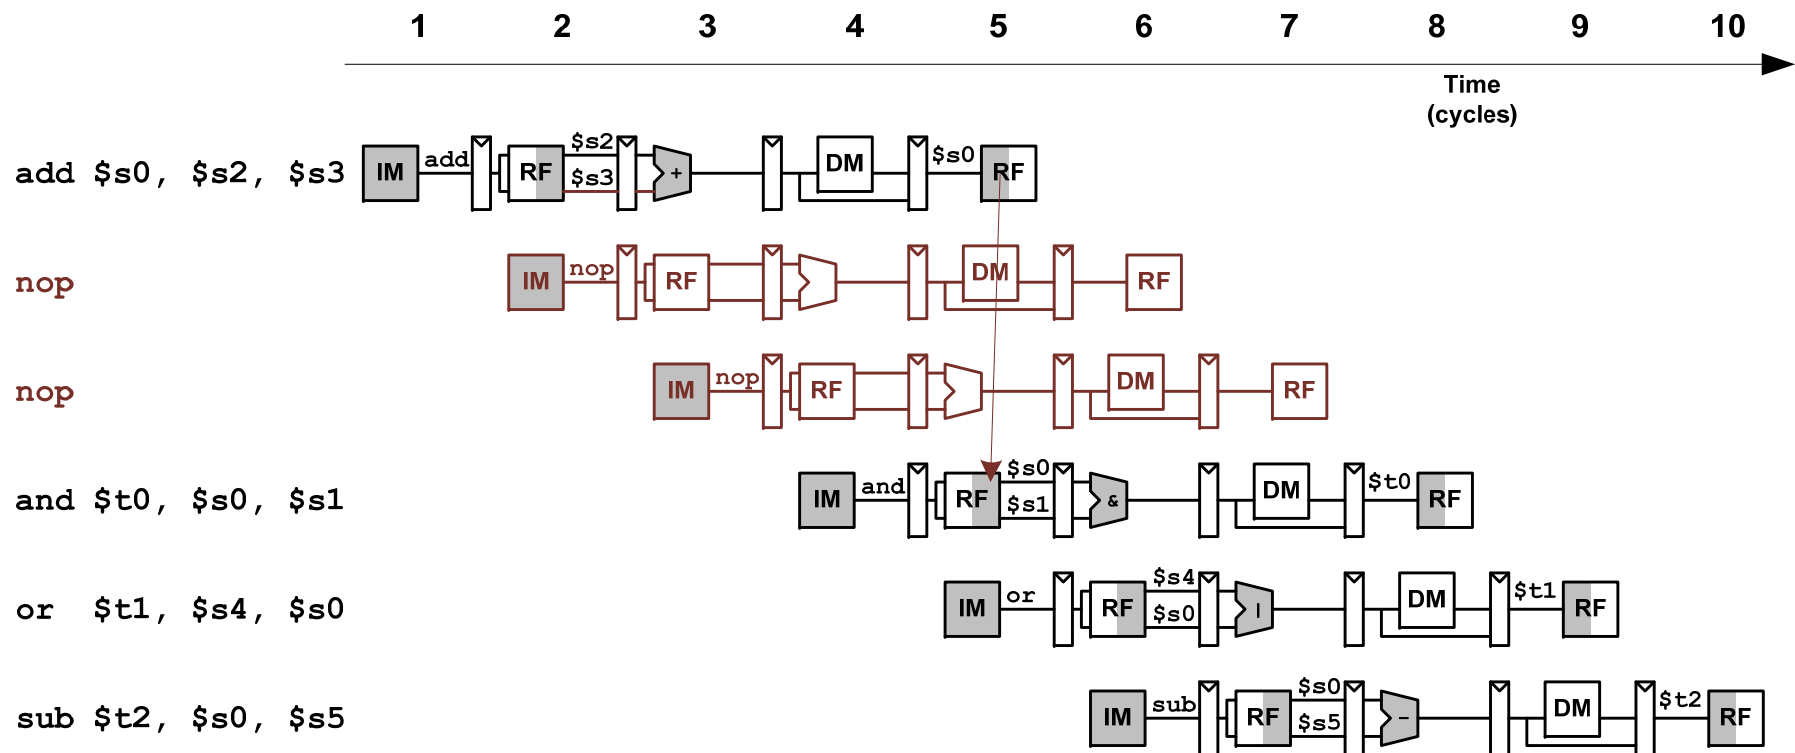
Handling Data Hazards /2

- **Handling Data Hazards:**
 1. Insert **nop** in code at the compile time
 2. Rearrange code at the compile time
 3. Forward data at run time
 4. Stall the processor at run time
- **Handle certain hazards through a special Hazard Unit with its own control signals**

Handling Data Hazards /3

■ Compile-Time Hazard Elimination:

- Insert enough **nop** for result to be ready
- No op(eration) or **nop** performs no operation
(**sll \$0, \$0, 0**)



Handling Data Hazards /4

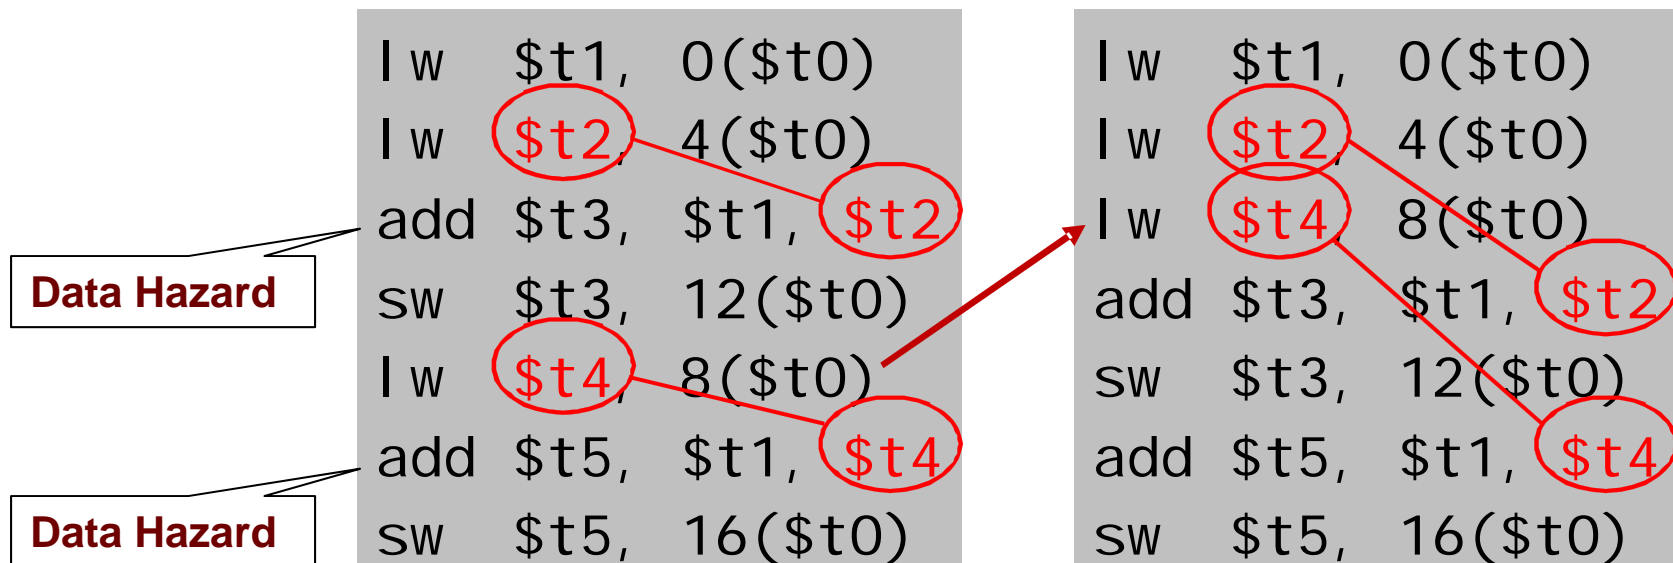
■ Compile-Time Hazard Elimination:

- Reorder code to avoid use of load result in the next instruction

- Example: MIPS code for

A = B + E;

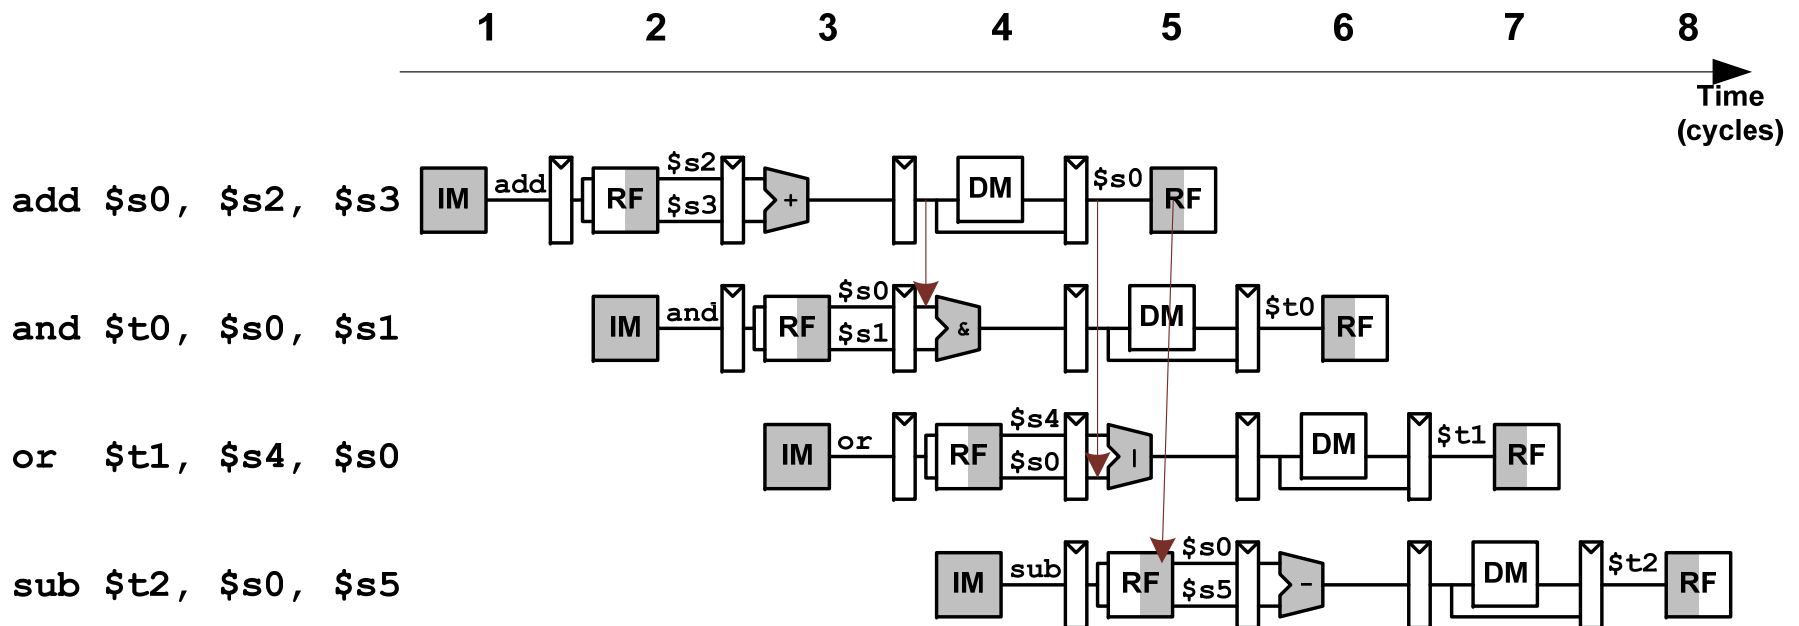
C = B + F;



Handling Data Hazards /5


■ Data Forwarding:

- Certain data hazards can be handled by forwarding (or bypassing) a result from the Memory or Writeback stage to a dependent instruction in the Execute stage
- Achieved by adding multiplexers in front of the ALU to select the operand from either the RF, the Memory or Writeback stage



Handling Data Hazards /6

■ Data Forwarding:



Hazard Unit

Handling Data Hazards /7

■ Data Forwarding via Hazard Unit:

- Forward to Execute stage from either:
Memory stage or Writeback stage

- Forwarding logic for ForwardAE (SrcA) :

```
if ((rsE != 0) AND (rsE == WriteRegM) AND RegWriteM)
    then ForwardAE = 10
else if ((rsE != 0) AND (rsE == WriteRegW) AND
    RegWriteW)
    then ForwardAE = 01
else ForwardAE = 00
```

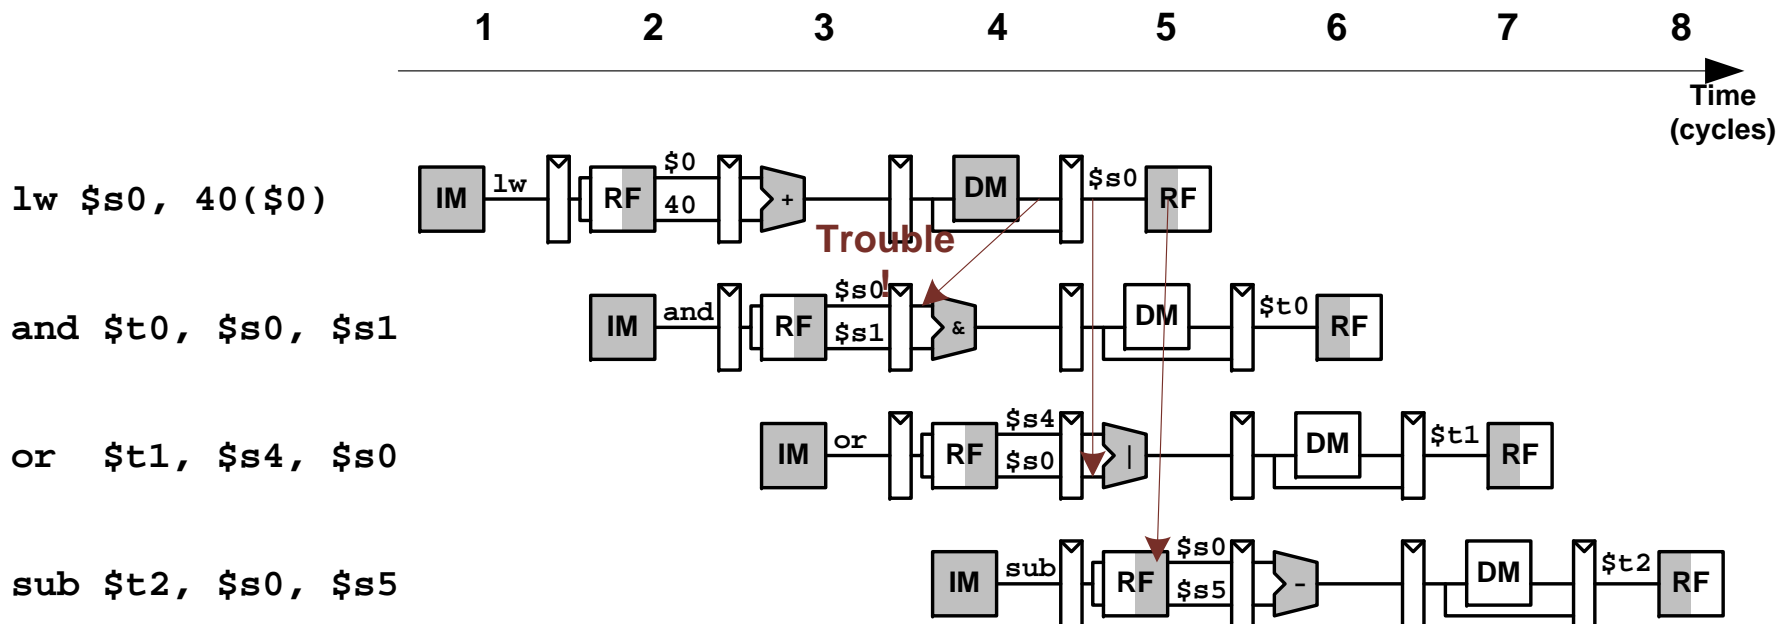
- Forwarding logic for ForwardBE (SrcA) :

```
if ((rtE != 0) AND (rtE == WriteRegM) AND RegWriteM)
    then ForwardBE = 10
else if ((rtE != 0) AND (rtE == WriteRegW) AND
    RegWriteW)
    then ForwardBE = 01
else ForwardBE = 00
```

Handling Data Hazards /8

■ Why stall at run time?

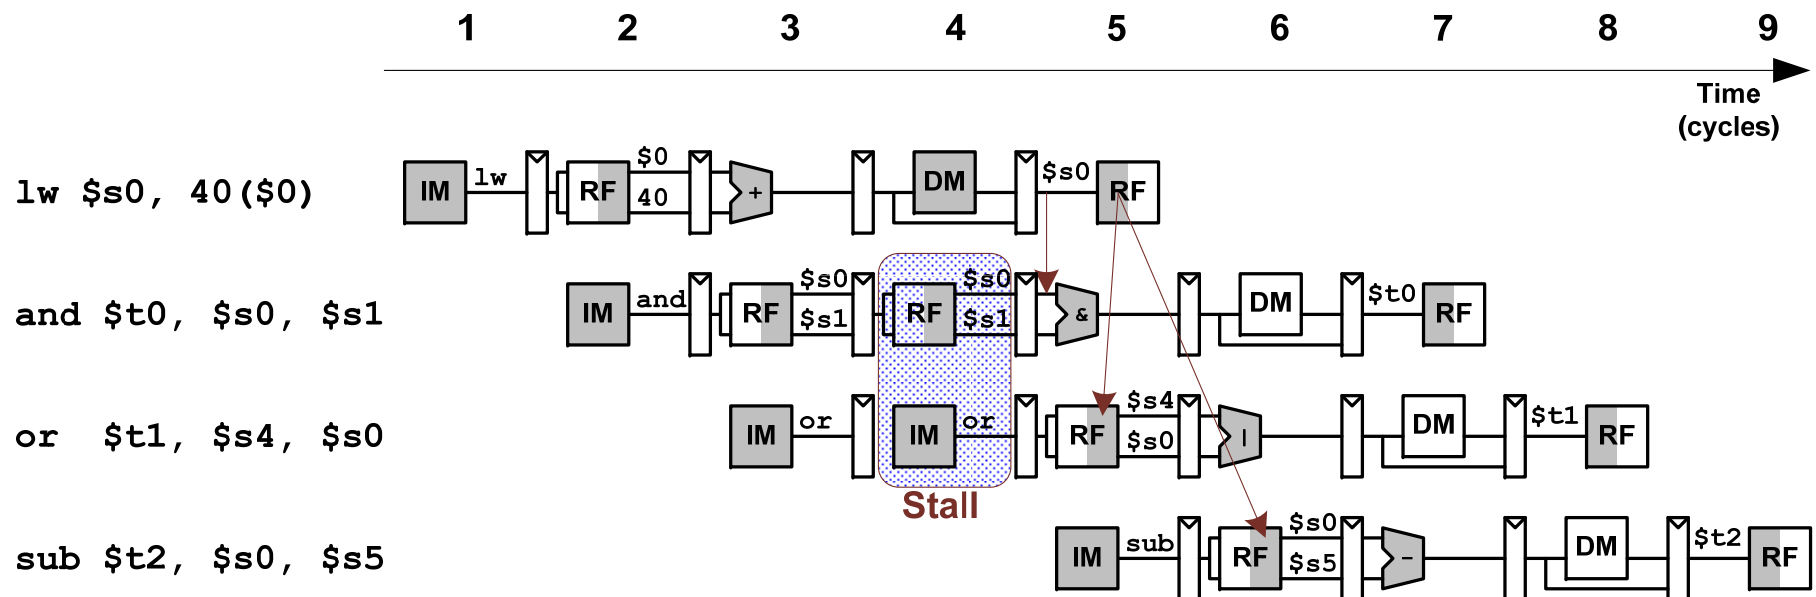
- The `lw` instruction experiences two-cycle latency
- `lw` does not finish reading data until the end of the Memory stage, so its result cannot be forwarded to the Execute stage



Handling Data Hazards /9

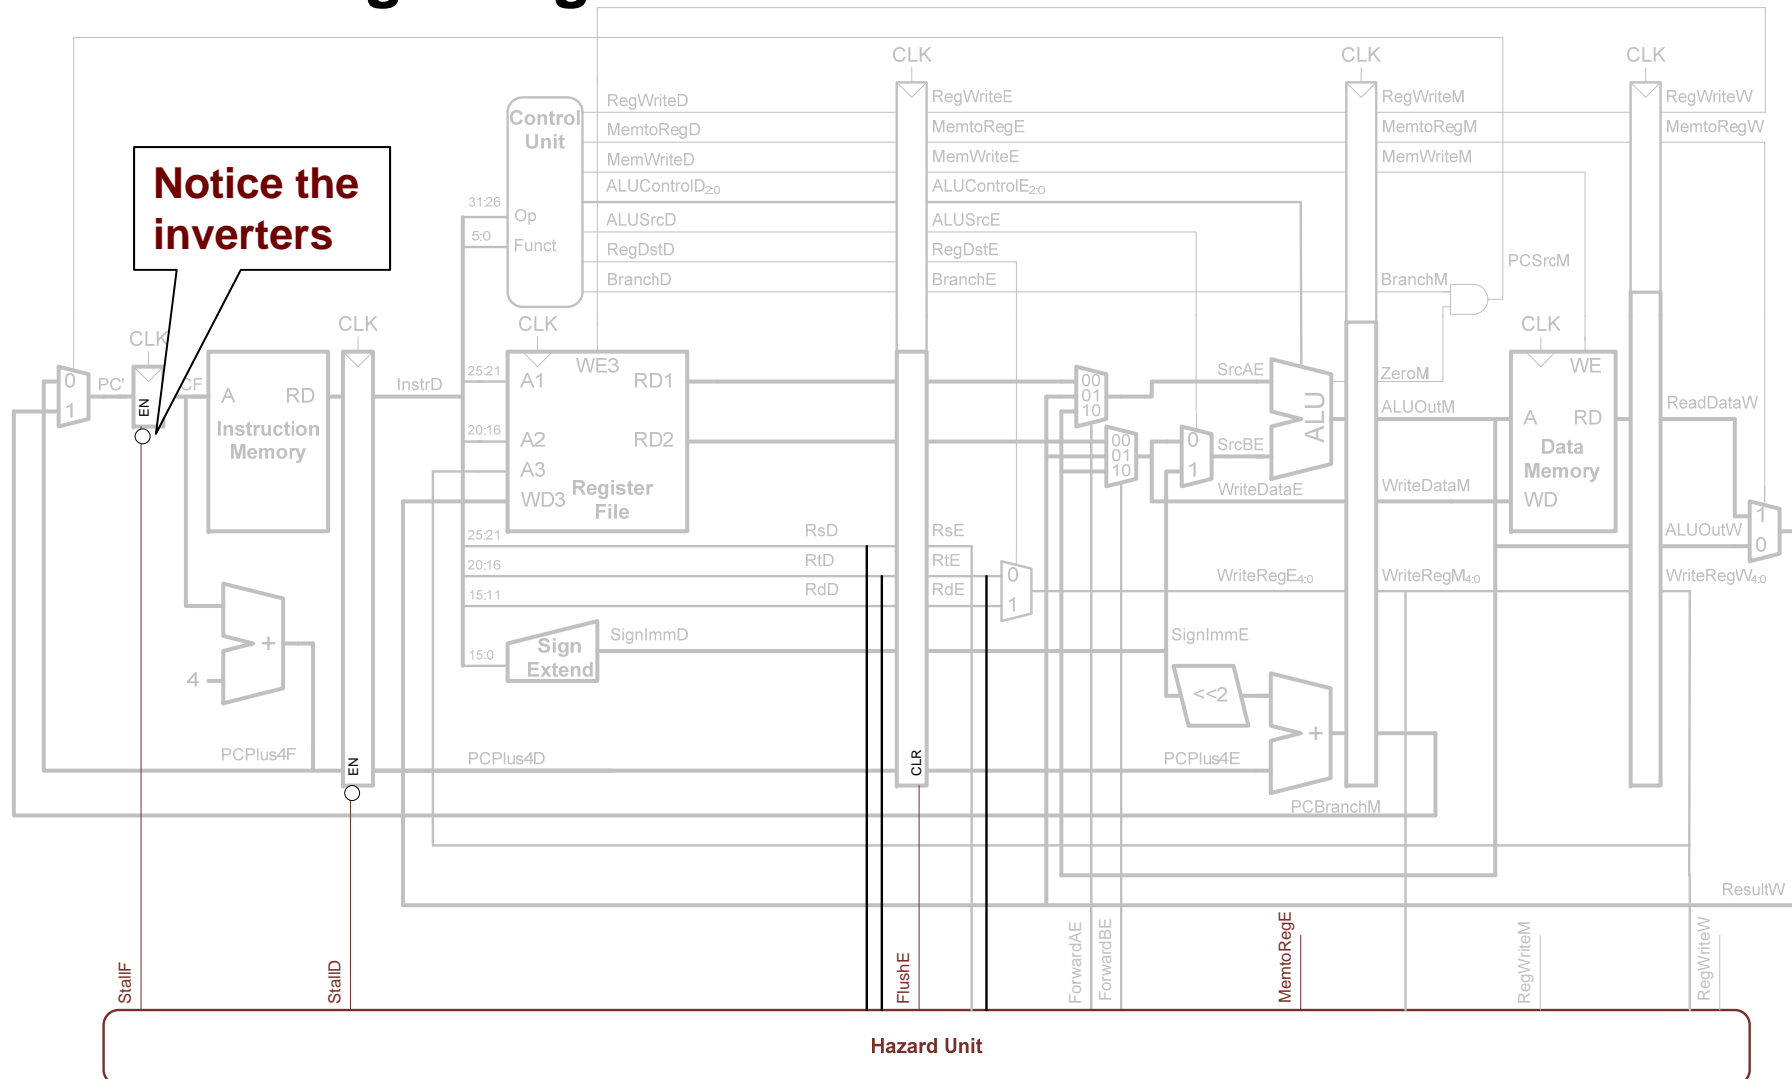
■ Stalling:

- **Solution: Hold up the pipeline (called stalling) until the needed data is available**
- The unused stage propagation through the pipeline is called a bubble, and it behaves like a `nop` instruction
- **The bubble is introduced by zeroing out the EX stage control signals during a Decode stall, so that the bubble performs no action and changes no architectural state**



Handling Data Hazards /10

■ Stalling using the Hazard Unit:



Handling Data Hazards /11

- **Stall by disabling the pipeline register, so that the contents do not change**
 - The pipeline register directly after the stalled stage must be cleared to prevent invalid information from moving forward
 - Stalls typically degrade the pipeline performance, so should be used sparingly
- **Stalling logic is relatively simple:**
 - Stall if dealing with lw and the destination register from EX is one of the operands in the ID stage
 - **Logic encoding of Stall bits that serve as EN bits:**
$$lwstall = ((rsD == rtE) \text{ OR } (rtD == rtE)) \text{ AND } MemtoRegE$$
$$StallF = StallD = FlushE = lwstall$$
 - **FlushE** is used to clear the contents of the EX stage register

Handling Control Hazards /1

■ **beq-related Control Hazards:**

- Branch value not determined until the fourth stage of pipeline
- Instructions after branch fetched before branch occurs
- **These instructions must be flushed if branch happens**

■ **Could we stall the pipeline?**

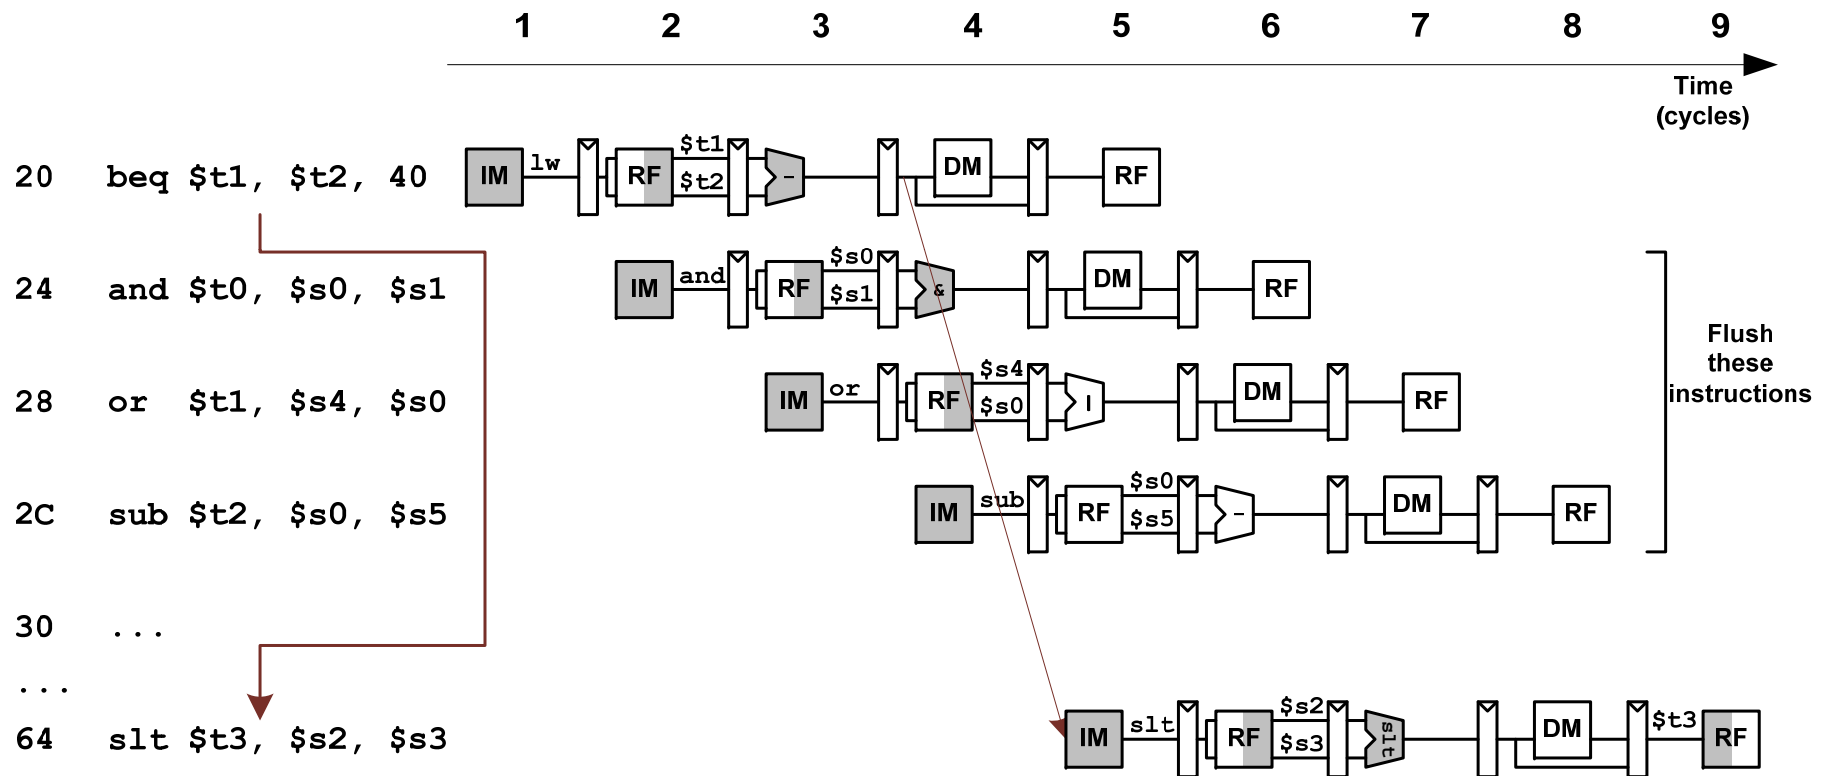
- Possibly yes, but we would have to stall for three cycles, which would be a significant performance penalty

■ **Instead, predict whether the branch will be taken**

- Start executing immediately based on the prediction
- If the prediction turns out to be wrong, throw out the loaded instructions (called instruction flush) and continue execution

Handling Control Hazards /2

- Instructions are flushed when the branch is taken:

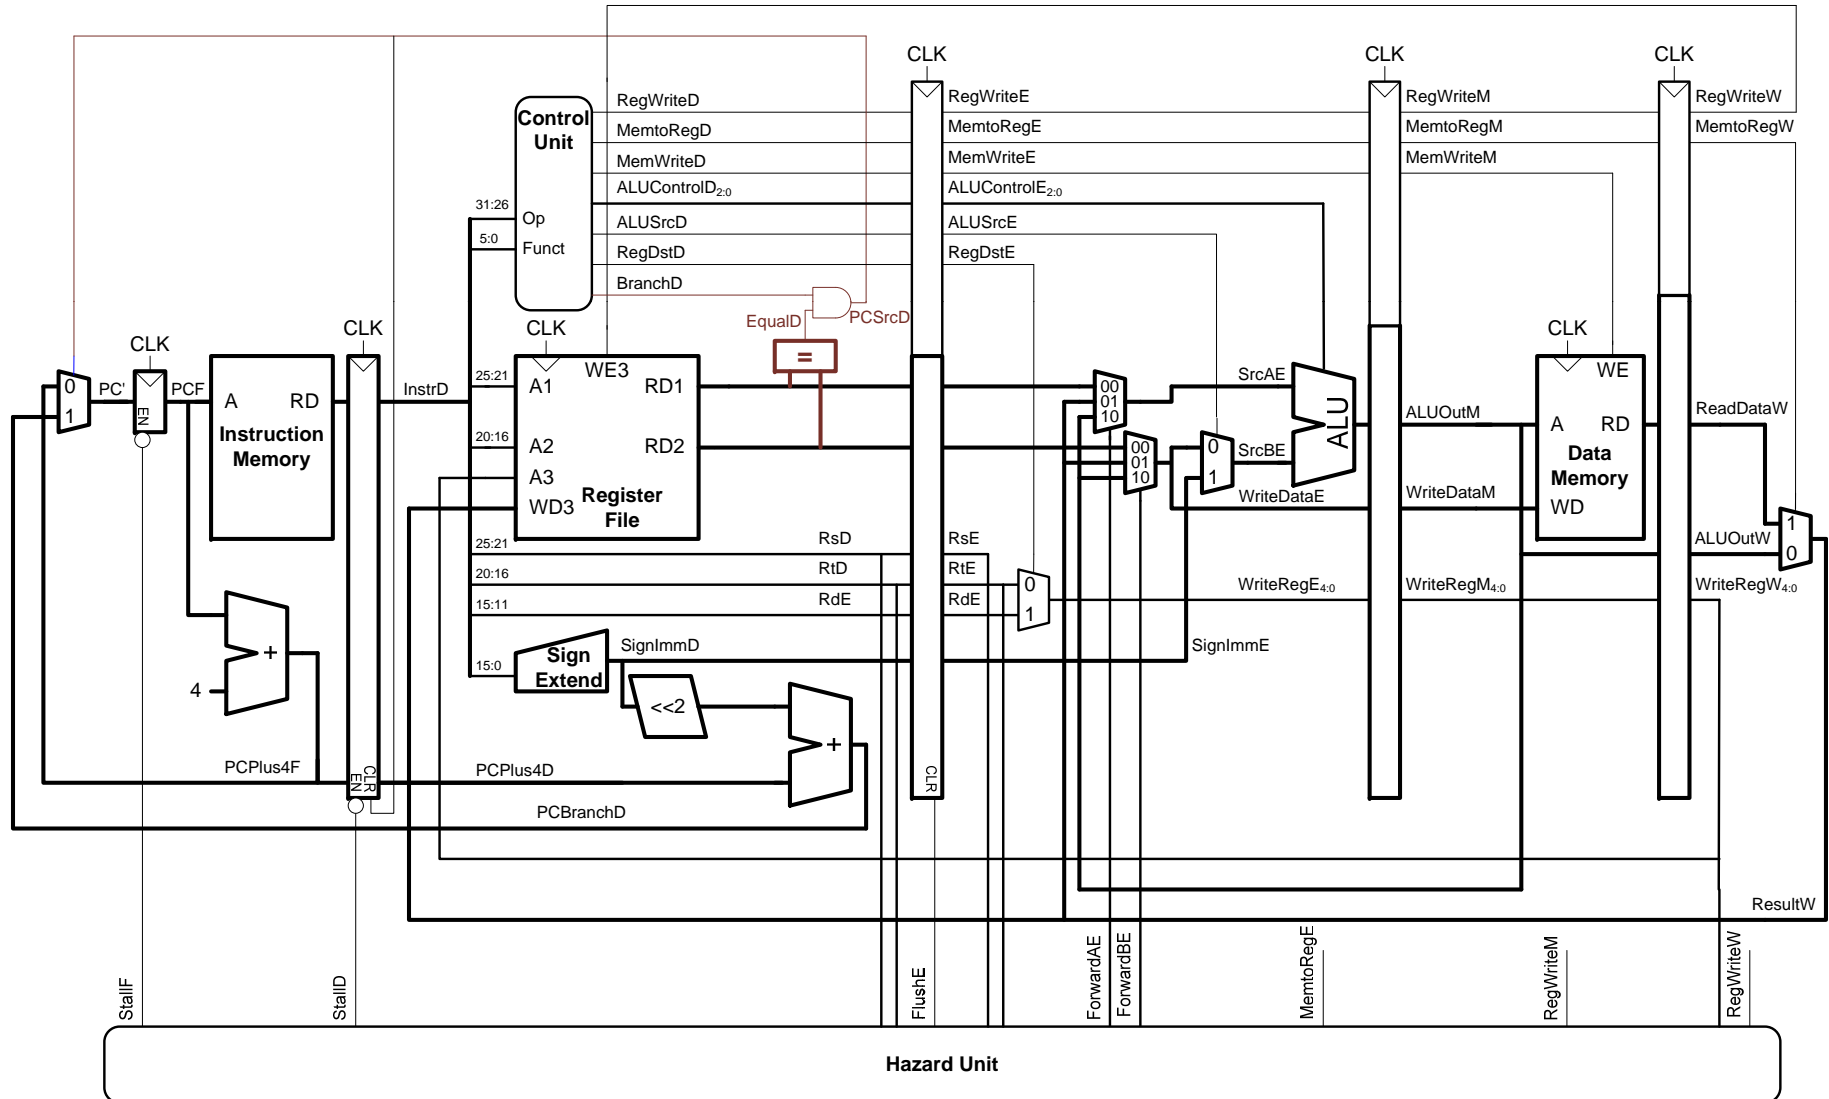


Handling Control Hazards /3

- **Branch Misprediction Penalty:**
 - Number of instruction flushed when branch is taken
 - May be reduced by determining the branch value earlier
- **Note that deciding whether the branch is taken simply requires an equality comparator**
 - This can be implemented much faster than using a subtraction and zero detection via ALU
 - Adding an equality comparator is also referred to as Early Branch Resolution
 - **By introducing the equality comparator into the ID stage, the branch misprediction penalty can be reduced to one (from three)**

Handling Control Hazards /4

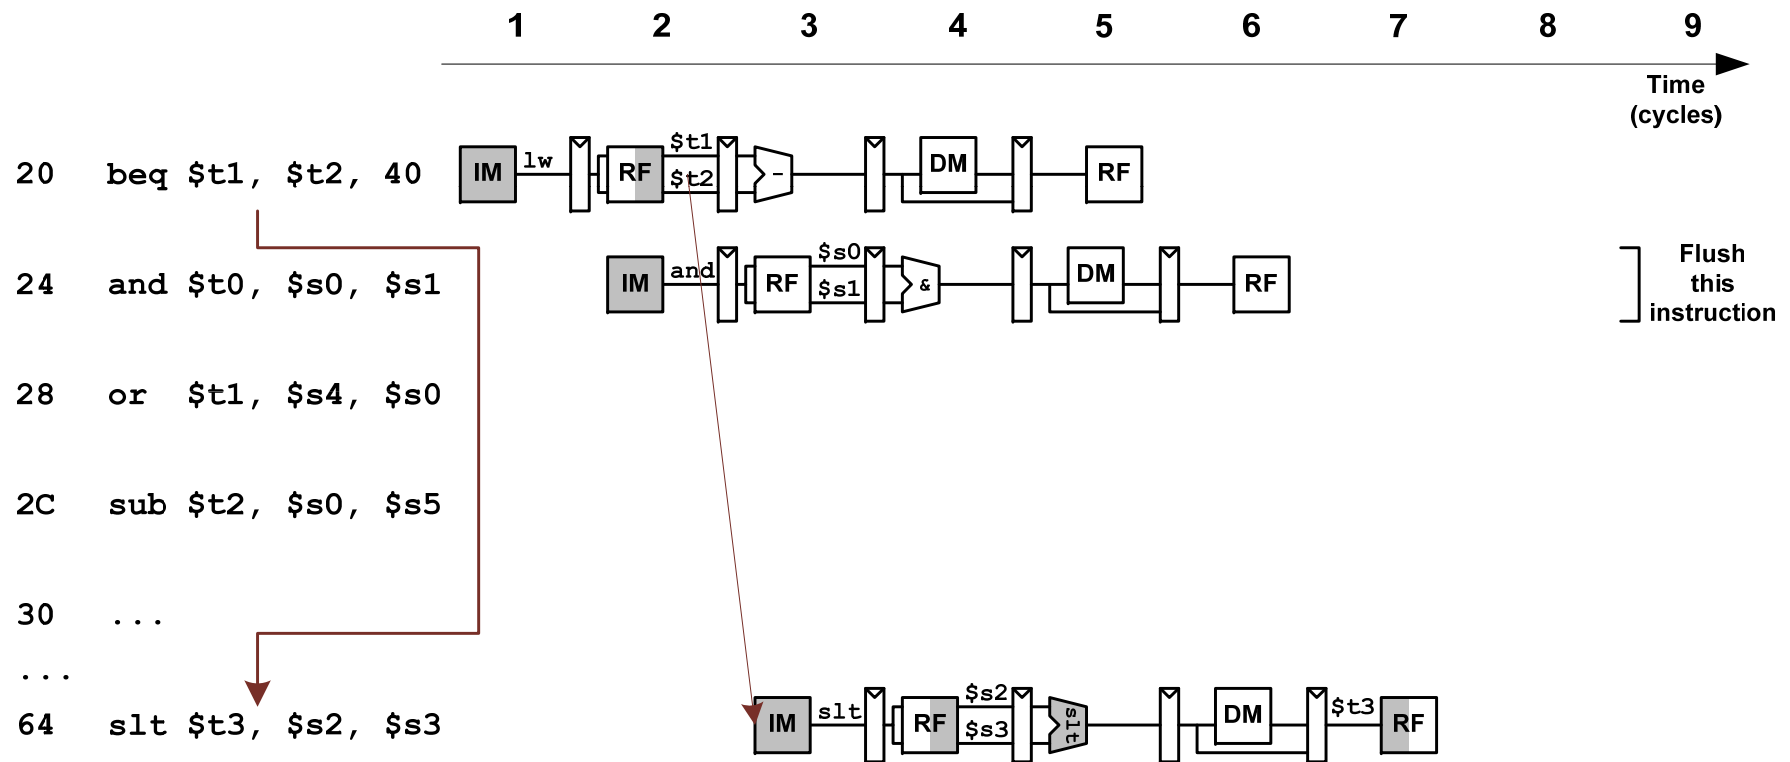
■ Early Branch Resolution:



Handling Control Hazards /5

■ Early Branch Resolution:

- Lowered the number of instructions that need to be flushed



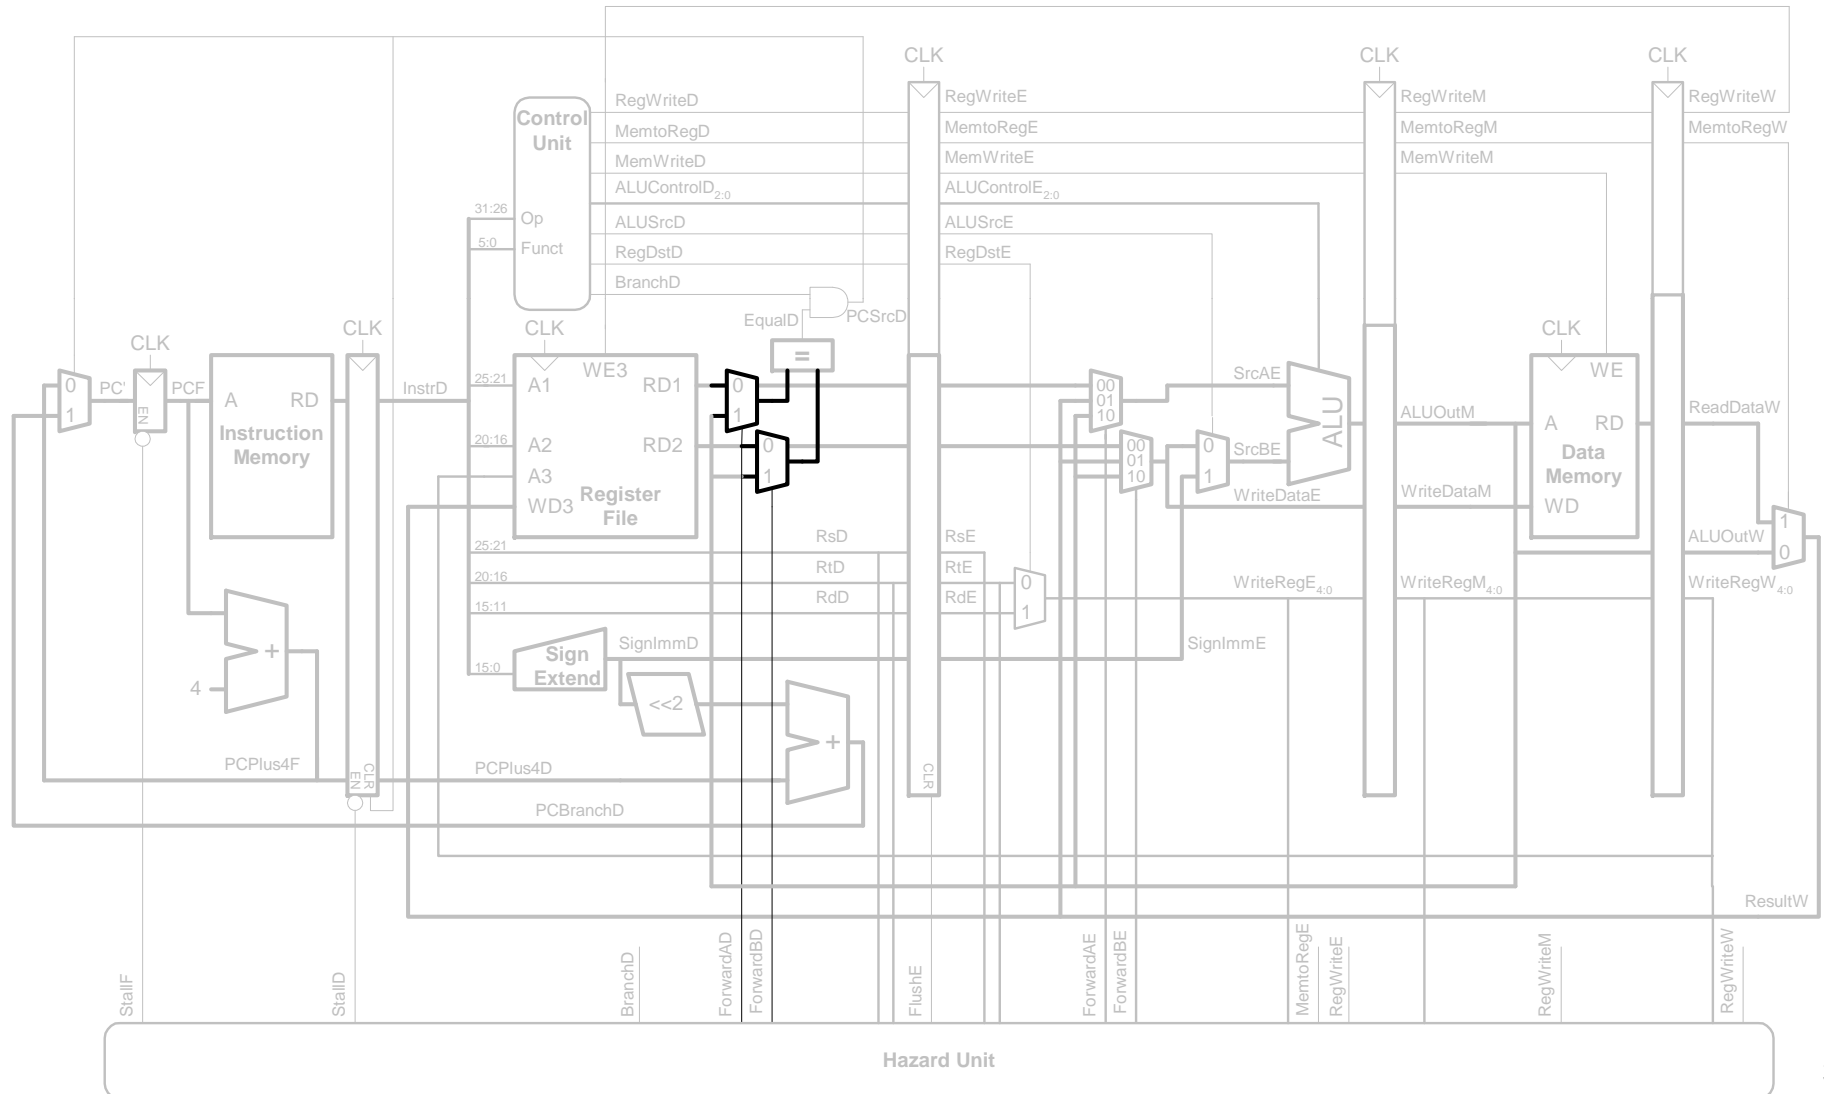
Handling Control Hazards /6

■ Handling Data and Control Hazards Together:

- Early branch resolution introduces a RAW hazard
- That is, if one of the source operands for the branching is being written by the previous instruction, it may not be ready yet
- For the WB stage, the hazard will not apply since the write will occur in the first part of the cycle and read in the second
- For the MEM stage of an ALU instruction, we could forward the needed value to the new equality comparator via multiplexors
 - Not using the memory unit, so we can freely forward the value
- For the EX stage of an ALU instruction or MEM stage of a lw instruction, we will need to delay (stall) the pipeline until the result is ready
 - Using the matching element in each option, so we need to stall
- **The special kind of stall that is introduced here is called the branch stall**

Handling Control Hazards /7

■ Handling Data and Control Hazards Together:



Handling Control Hazards /8

■ Handling Data and Control Hazards Together:

■ Forwarding Logic:

`ForwardAD = (rsD != 0) AND (rsD == WriteRegM) AND
RegWriteM`

`ForwardBD = (rtD != 0) AND (rtD == WriteRegM) AND
RegWriteM`

■ Stalling Logic:

`branchstall = (BranchD AND RegWriteE AND
(WriteRegE == rsD OR WriteRegE == rtD))`

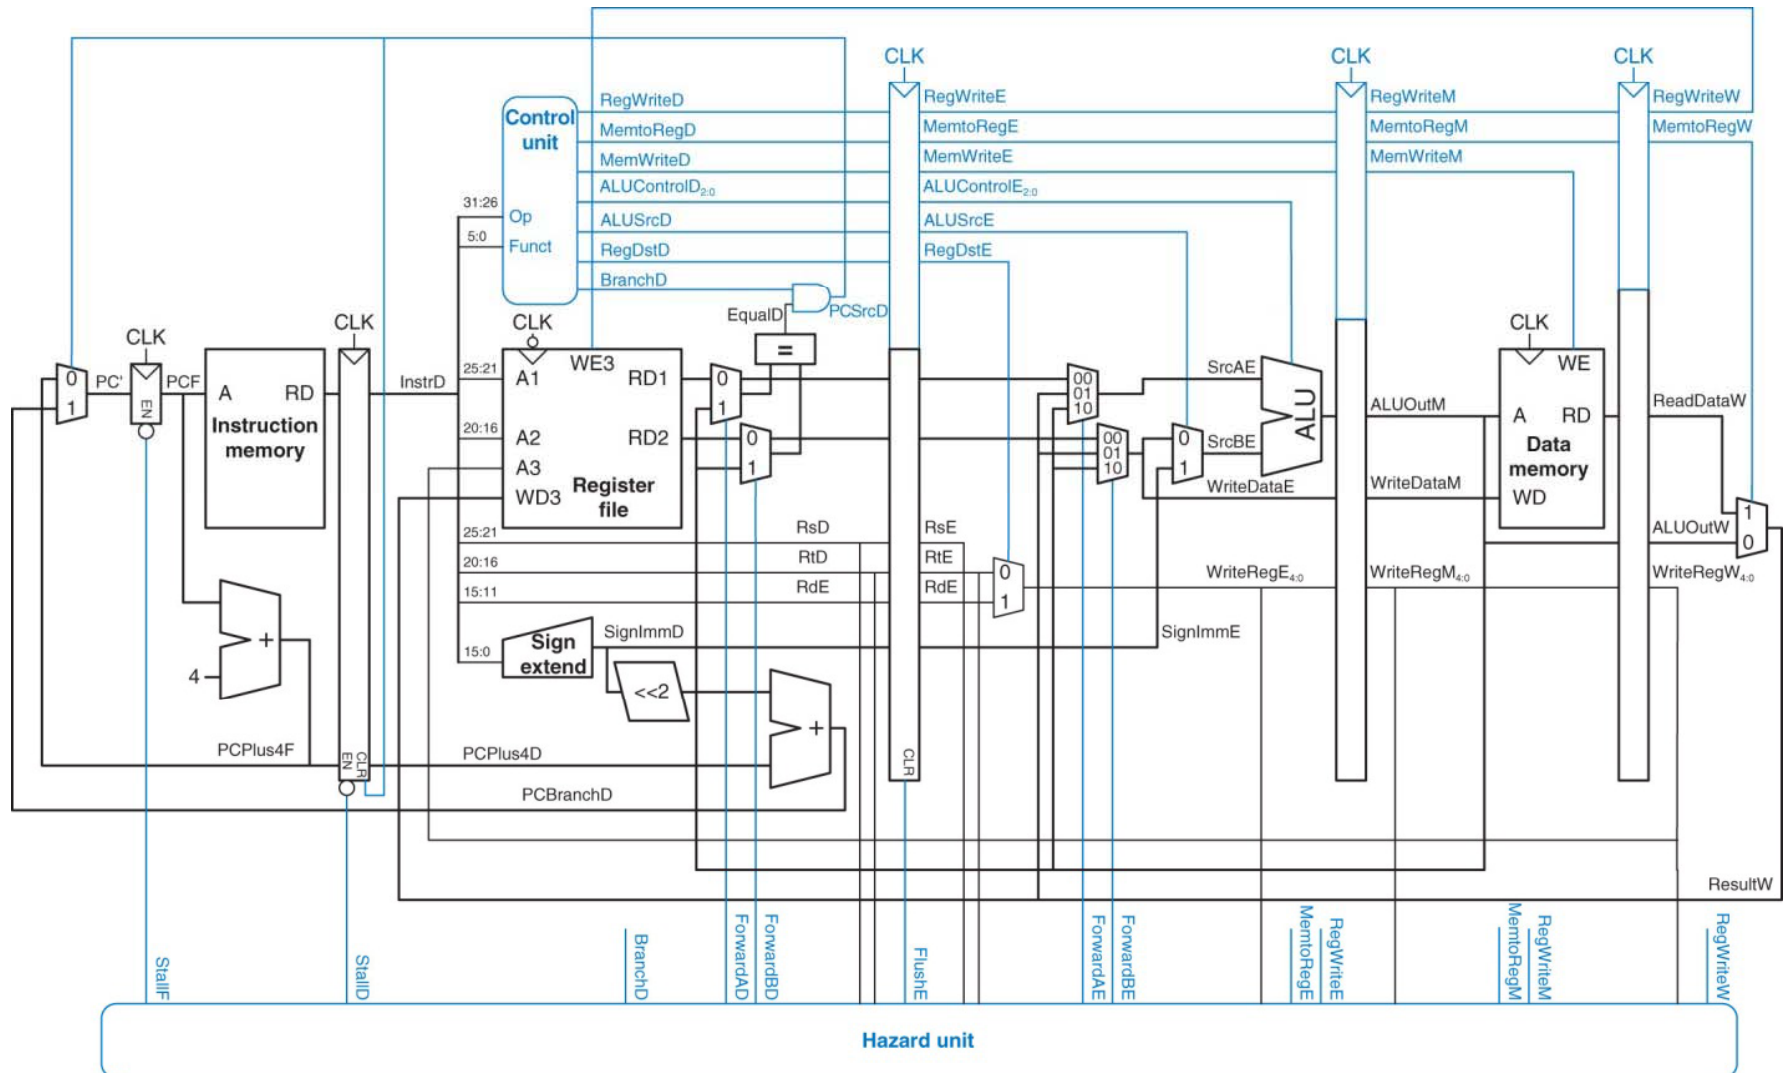
`OR`

`(BranchD AND MemtoRegM AND
(WriteRegM == rsD OR WriteRegM == rtD))`

`StallF = StallD = FlushE = lwstall OR branchstall`

Handling Control Hazards /9

■ Final Pipelined MIPS Datapath Design:



Handling Control Hazards /10

■ **Branch prediction to improve average CPI:**

- Ideal pipelined processor: $CPI = 1$
- Branch misprediction increases the average CPI

■ **Static branch prediction:**

- Check direction of branch (forward or backward)
- If backward, predict taken
- Else, predict not taken

■ **Dynamic branch prediction:**

- **Keep the history of the last several hundred branches in the branch target buffer, and record:**
 - Branch destination
 - Whether the branch was taken
- Good prediction reduces fraction of branches requiring a flush

Handling Control Hazards /11

■ Branch Prediction Example:

```
add  $s1, $0, $0      # sum = 0
add  $s0, $0, $0      # i   = 0
addi $t0, $0, 10      # $t0 = 10
for:
    beq  $s0, $t0, done  # if i == 10, branch
    add  $s1, $s1, $s0    # sum = sum + i
    addi $s0, $s0, 1      # increment i
    j    for
done:
```

The loop
continues
until i == 10


■ Simple one-bit dynamic branch predictor:

- Check if the last branch was taken
- Repeat the same action from the previous cycle
- **One-bit predictor is simple to design, but it mispredicts the first and the last branch of each loop**

Handling Control Hazards /12

■ Two-bit branch predictor:

- Introduces four states: strongly taken, weakly taken, weakly not taken, and strongly not taken
- Changes the next action after two mispredictions, or two correct predictions, respectively
- **Two-bit predictor mispredicts the last branch of a loop**



Pipelined Performance Example /1

■ Example:

- 40% of loads used by the next instruction
- 25% of branches mispredicted
- All jumps flush the next instruction
- Also, 25% loads, 10% stores, 11% branches, 2% jumps, and 52% R-type instructions are executed

■ What is the average CPI?

- Load/Branch CPI = 1 when no stalling, 2 when stalling
- $CPI_{lw} = 1(0.6) + 2(0.4) = 1.4$
- $CPI_{beq} = 1(0.75) + 2(0.25) = 1.25$
- Average CPI =
 $(0.25)(1.4) + (0.1)(1) + (0.11)(1.25) + (0.02)(2) + (0.52)(1) =$
Average CPI = 1.15

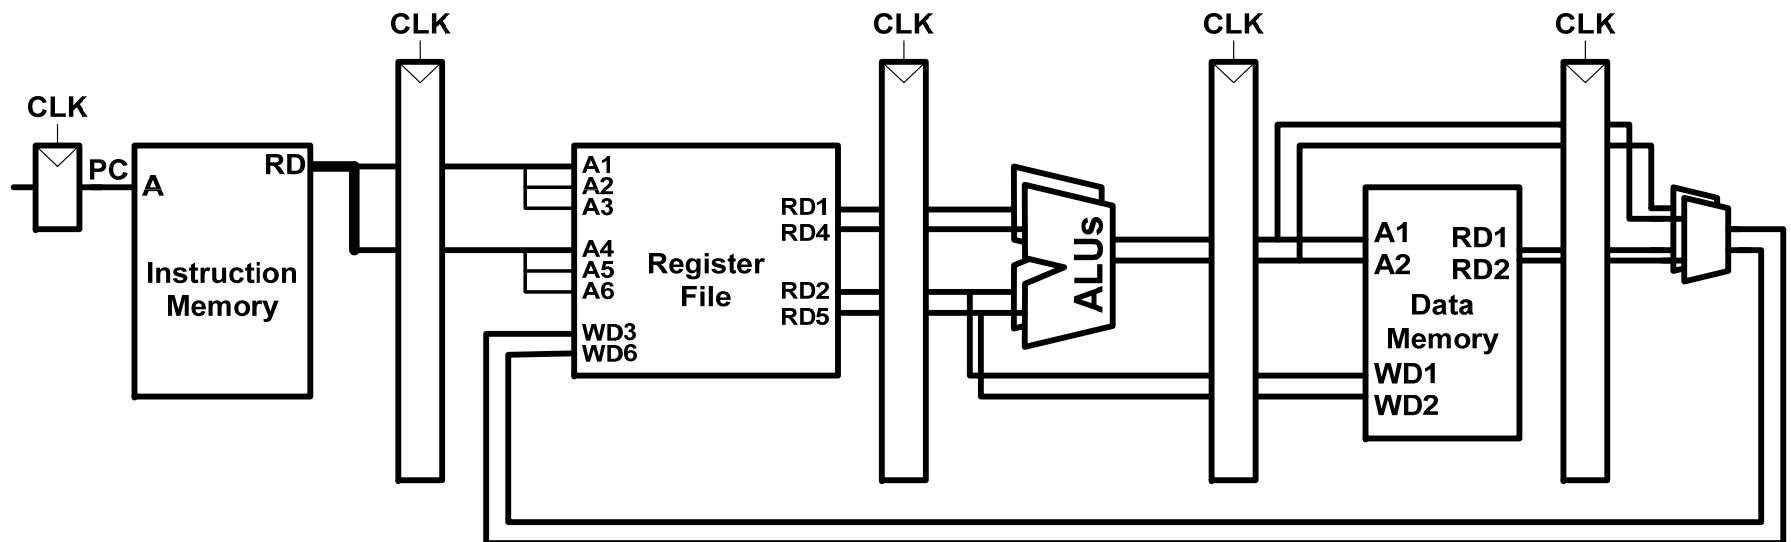
Pipelined Performance Example /2

- **Sample values for the clock-cycle times (T_c):**
 - Single-cycle processor T_c : 925 ps
 - Single-cycle processor CPI: 1
 - Multi-cycle processor T_c : 325 ps
 - Multi-cycle processor CPI: 4.04
 - **Pipelined processor T_c : 550ps**
 - **Pipelined processor CPI: 1.15**
- **So which of the three processor designs will perform faster when executing 100 billion gcc instructions?**
 - Single-cycle processor execution time = 92.5 seconds
 - Multi-cycle processor execution time = 131.3 seconds
 - Pipelined processor execution time =
of instructions \times CPI $\times T_c = (100 \times 10^9) \times 1.15 \times 550 \times 10^{-12}$
= 63.25 seconds

Advanced MIPS Pipeline Design /1

■ Superscalar Datapath:

- Multiple copies of the pipeline datapath execute multiple instructions at once
- The datapath fetches two instructions at one time
- It contains a register file with six input register ports, two ALUs, and a data memory unit with two input address ports
- **Dependencies between instructions may make it tricky to issue multiple instructions at once**

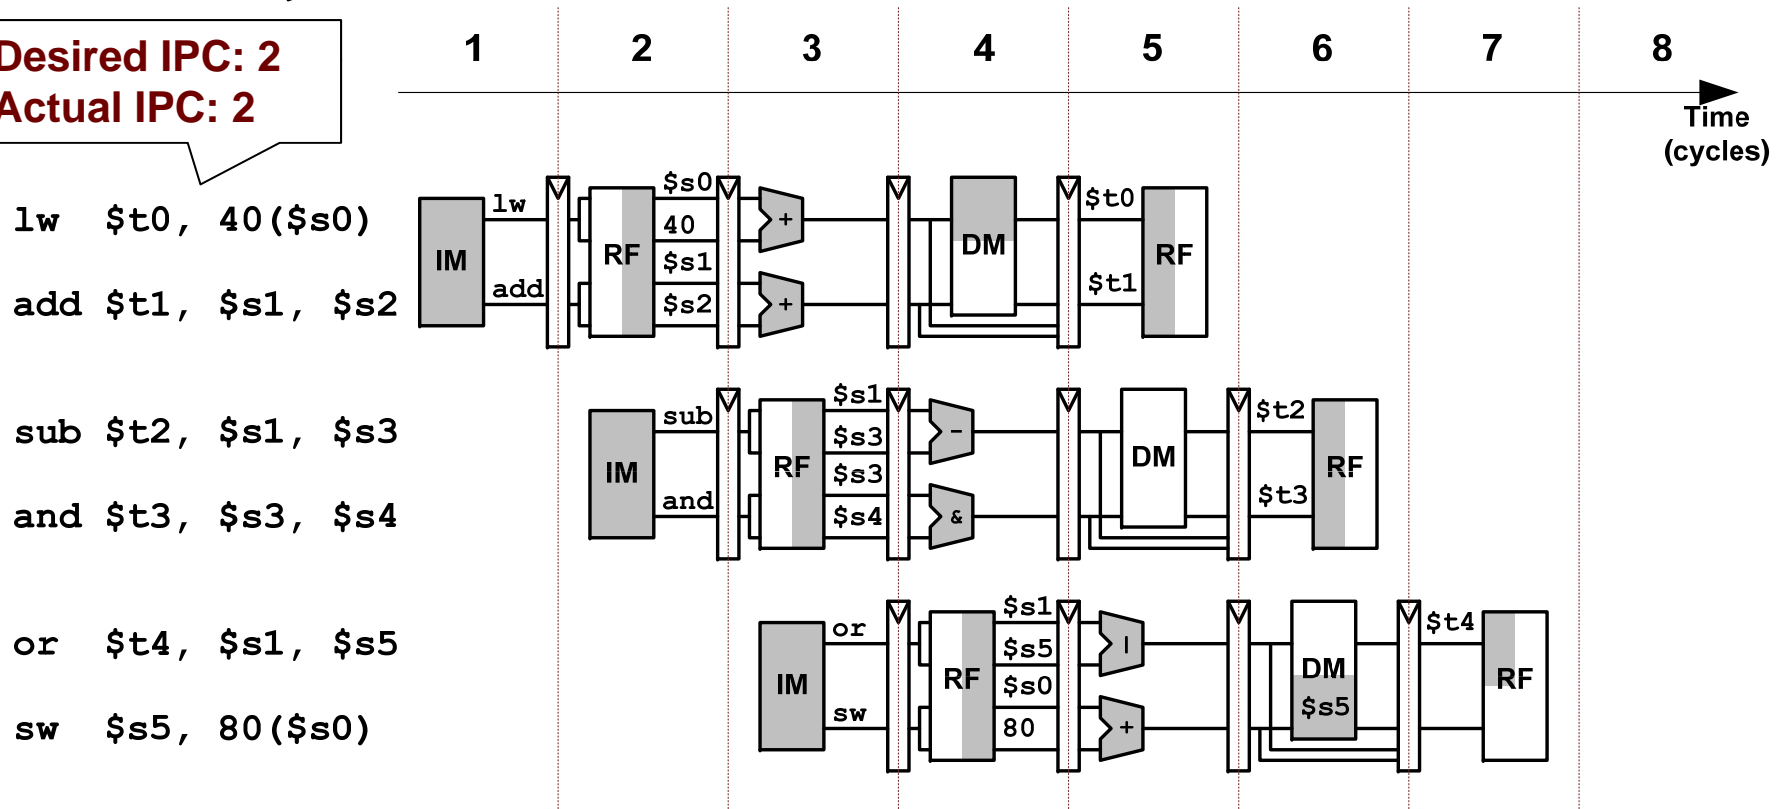


Advanced MIPS Pipeline Design /2

■ Superscalar Datapath Example:

IPC value refers to the number of instructions issued per cycle

Desired IPC: 2
Actual IPC: 2



Advanced MIPS Pipeline Design /3

■ Superscalar Datapath Example with Dependencies:

Desired IPC: 2

Actual IPC:
 $6/5 = 1.17$

lw \$t0, 40(\$s0)

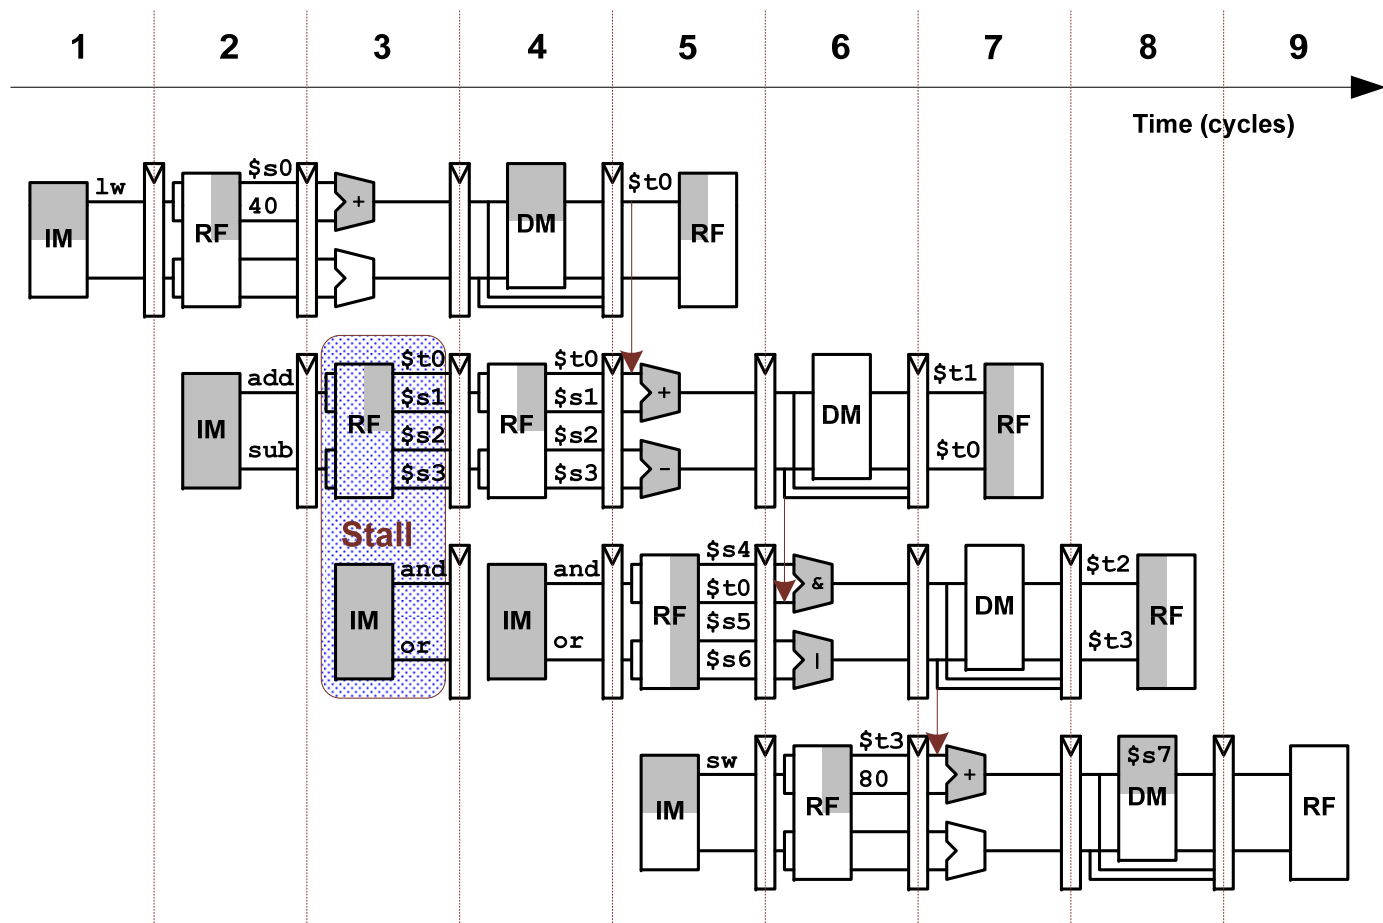
add \$t1, \$t0, \$s1

sub \$t0, \$s2, \$s3

and \$t2, \$s4, \$t0

or \$t3, \$s5, \$s6

sw \$s7, 80(\$t3)



Advanced MIPS Pipeline Design /4

■ **Deep Pipelining:**

- Split the pipeline into more stages
- Each stage requires less logic to run, so it runs faster
- It is typical that 10 to 20 stages are used
 - We have covered the classical 5-stage MIPS pipeline
- The maximum number of pipelines is limited by additional dependencies between instructions, sequencing overhead, increased power requirements, and increased cost

■ **Additional stages introduce additional hazards**

- Increasing the number of stages increases the number of additional dependencies between instructions, which can lead to additional data and control hazards
- Some hazards can be resolved by forwarding, but some may require stalls, which increases the average CPI

Advanced MIPS Pipeline Design /5

■ **Out-of-Order Processor:**

- Looks ahead across multiple instructions
- Issues as many instructions as possible at once
- Issues instructions out of order as long as all dependencies are not violated

■ **Dependencies to consider:**

- **RAW (read after write):** One instruction writes, later instruction reads a register before the first instruction finished writing it
 - Occurs in the classical MIPS pipeline; essential hazard
- **WAR (write after read):** One instruction reads, later instruction writes over a value before the intended instruction has read it
 - Also called anti-dependence; non-essential hazard
- **WAW (write after write):** One instruction writes, later instruction is moved forward so it writes over the written value
 - Also called output dependence; non-essential hazard

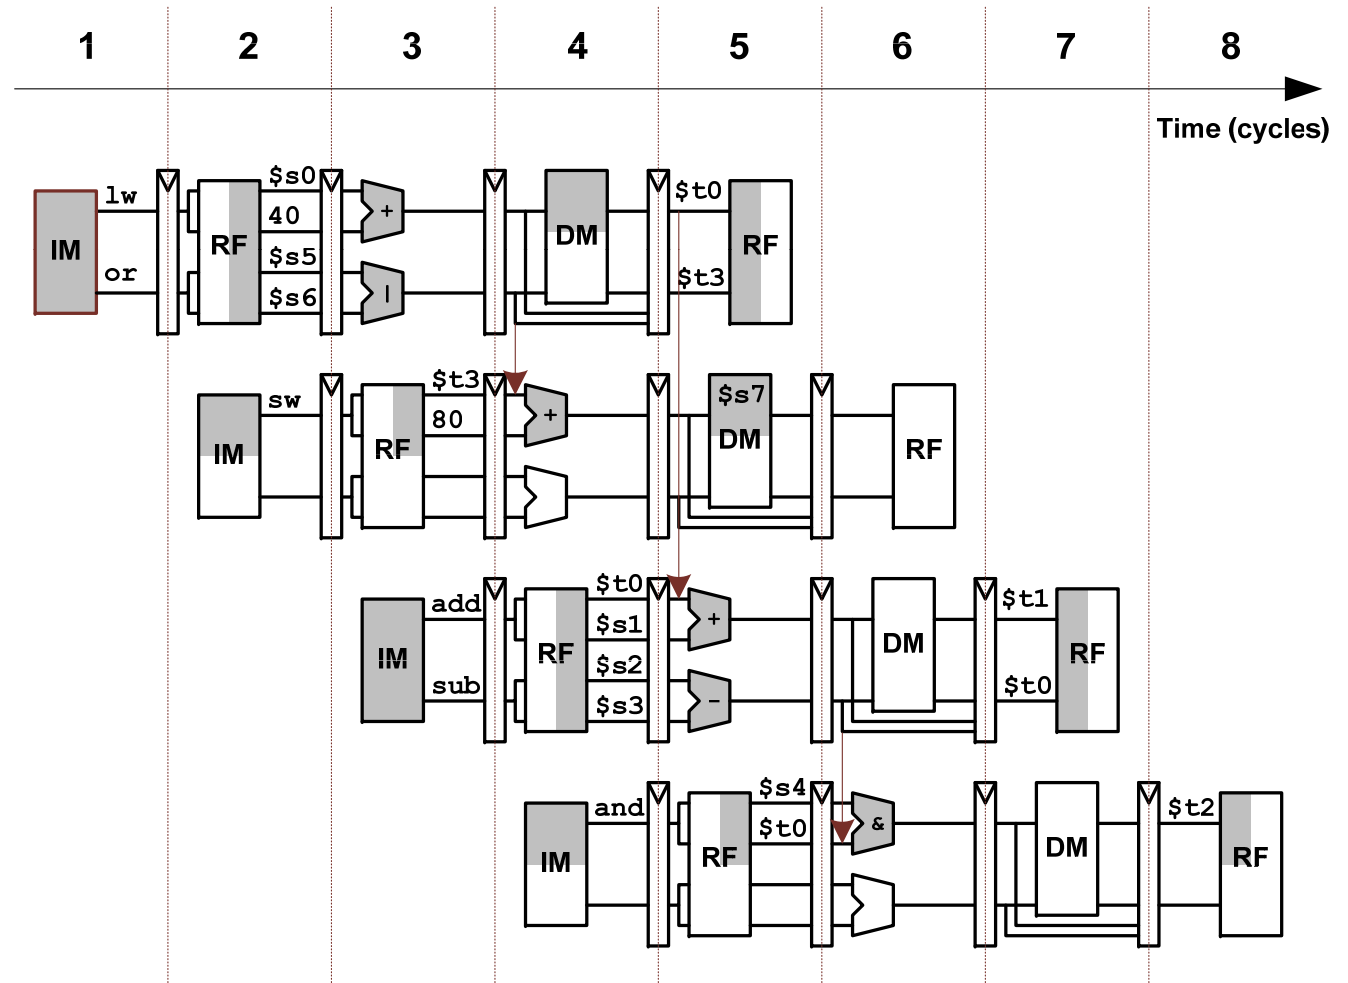
Advanced MIPS Pipeline Design /6

■ Out-of-Order Processor Example:

Desired IPC: 2
Actual IPC: $6/4 = 1.5$

`lw $t0, 40($s0)`
`or $t3, $s5, $s6`
`sw $s7, 80($t3)`
`add $t1, $t0, $s1`
`sub $t0, $s2, $s3`
`and $t2, $s4, $t0`

RAW
 two cycle latency between load and use of \$t0
 RAW
 WAR
 RAW



Advanced MIPS Pipeline Design /7

■ **Out-of-Order Processor Implementation:**

■ **Use a table (called scoreboard) that keeps track of:**

- Instructions waiting to issue
- Available functional units
- Dependencies

■ The processor executes the maximum number of instructions based on the current status of the scoreboard

■ **Instruction level parallelism (ILP):**

- Theoretical ILP can be quite large for architectures with many architectural elements and very accurate branch predictors
- **In practice however, the number of instructions that can be issued simultaneously is on average less than 3**

Advanced MIPS Pipeline Design /8

■ Out-of-Order Processor Register Renaming:

- Add renaming registers to the datapath
- For instance, the datapath may include additional renaming registers called \$rn0 - \$rn19
- **These registers cannot be used directly, but can be used to resolve hazards**
- For instance, if \$t0 was used in two conflicting instructions, \$t0 could be renamed to \$rn0 in one of the two instructions

■ Register Renaming Example:

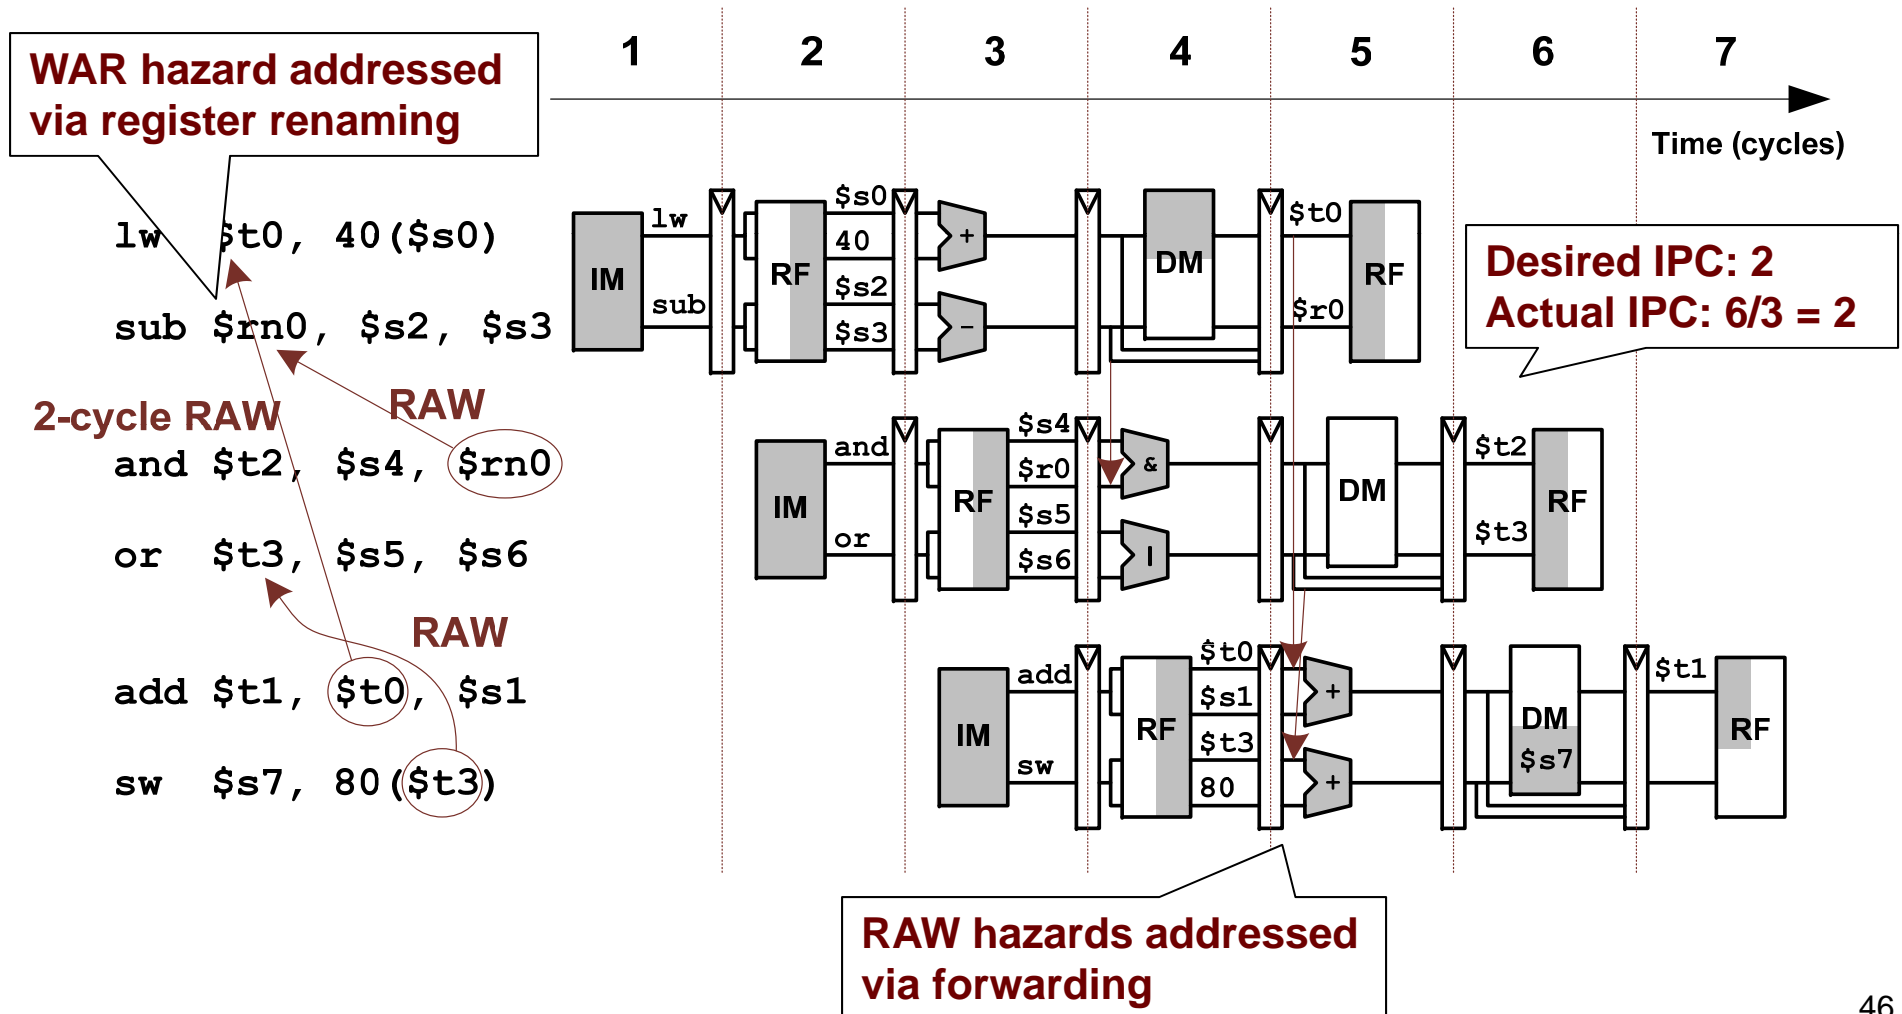
```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```



**WAR hazard may
be introduced**

Advanced MIPS Pipeline Design /9

■ Out-of-Order Processor Register Renaming:



Advanced MIPS Pipeline Design /10

■ **Process:**

- A program running on a computer
- Multiple processes can run at once
- For example, surfing Web, playing music, writing a paper

■ **Thread:**

- A unit of program execution
- Each process has multiple threads
- For example, a word processor may have threads for typing, spell checking, printing

■ **Multiprocessors:**

- Multiple processors (cores) with a method for communication between them (more on this later in the course)

Advanced MIPS Pipeline Design /11

■ **Threads in a conventional processor:**

- One thread runs at one time
- When one thread stalls (for example, waiting for memory):
 - Architectural state of that thread is stored
 - Architectural state of a waiting thread is loaded into the processor and that thread is then run
 - **This process is called context switching**
- Appears to a user like threads are running simultaneously

Advanced MIPS Pipeline Design /12

■ **Multithreading:**

- Multiple copies of the architectural state
- Multiple threads active at once:
 - When one thread stalls, another runs immediately
 - If one thread cannot keep all execution units busy, another thread can use them
- **Does not increase instruction-level parallelism (ILP) of a single thread, but it does increase throughput**
- Intel refers to this process as “hyper-threading”

Advanced MIPS Pipeline Design /13


■ Single Instruction Multiple Data (SIMD):

- One SIMD instruction acts on multiple pieces of data at once
 - Typically implemented as a coprocessor (e.g., ARM core)
- Common application: graphics
- Perform short arithmetic operations
 - Also called packed arithmetic

■ For example, add four 8-bit elements:

- Modify the ALU to ignore carry-outs between packets

padd8 \$s2, \$s0, \$s1



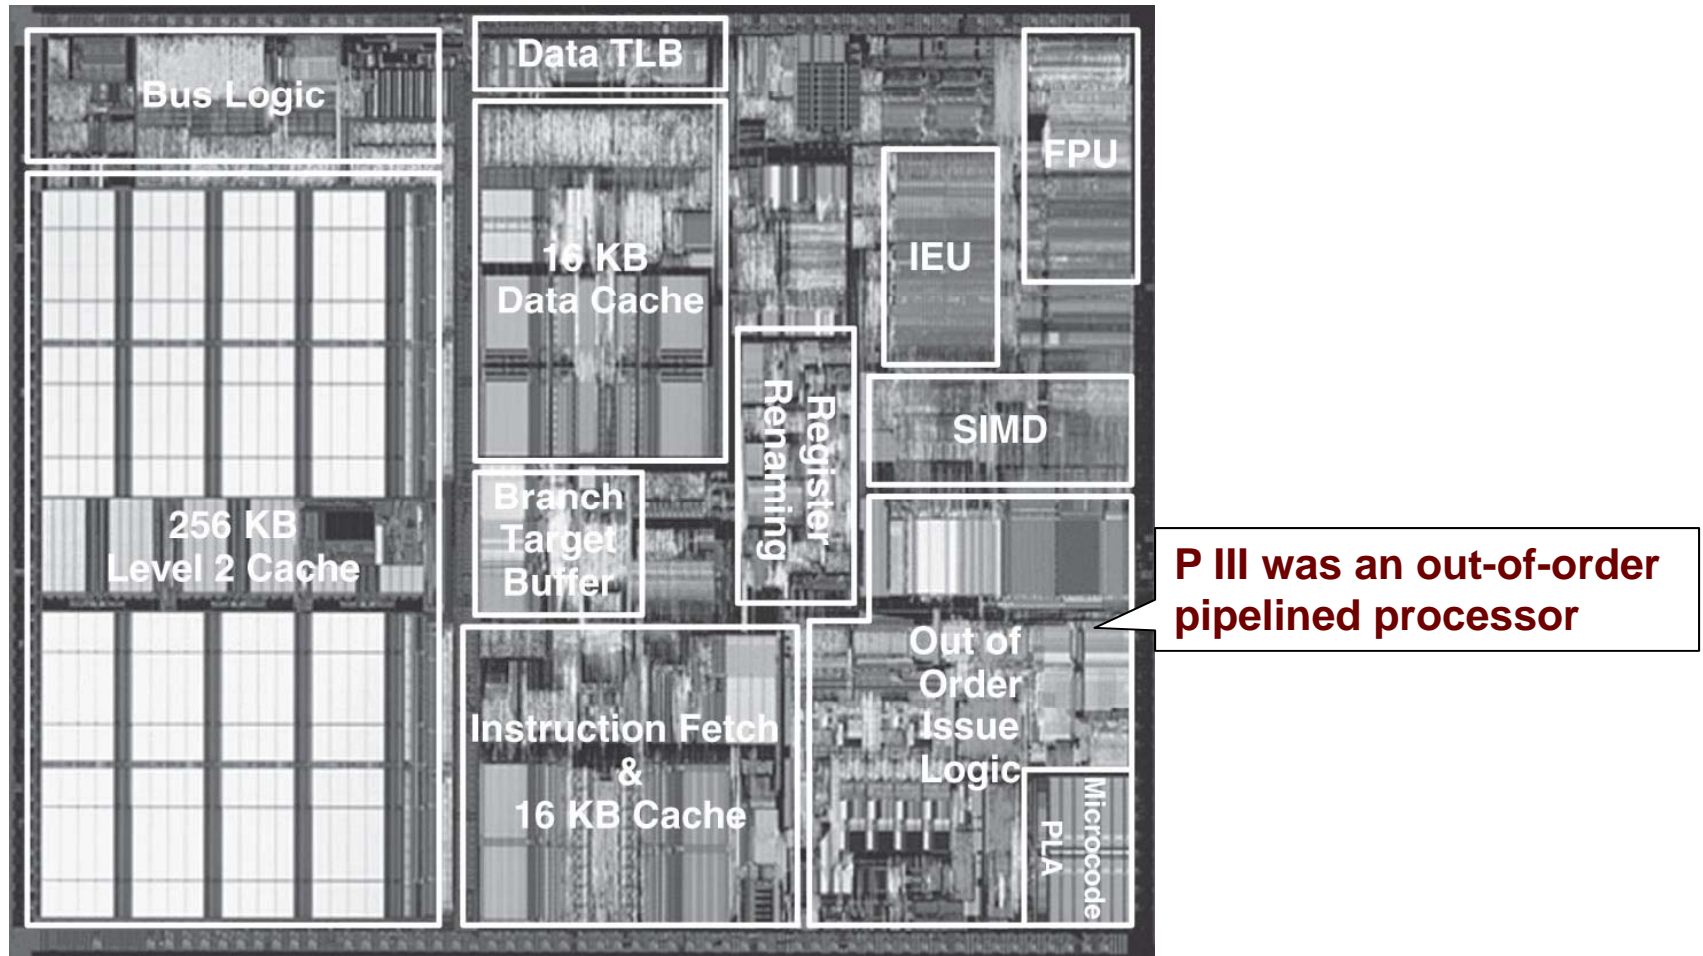
Advanced MIPS Pipeline Design /14

■ **More on Single Instruction Multiple Data (SIMD):**

- Quite useful when dealing with data elements that are smaller than 32 bits
- For instance, a pixel may store 8-bit information about a red, green, or blue component at one time
- Using an entire 32-bit value to perform 8-bit arithmetic would be wasteful (the upper 24 bits would go unused)
- **Instead, package four 8-bit packets into a 32-bit word, and perform arithmetic on multiple values at one time**
- Double-value words (64-bit words) can be even more helpful since those can be used to package eight 8-bit data elements, or four 16-bit data elements
- SIMD instructions can also be used in floating-point arithmetic

Visualizing Computer Architectures /1

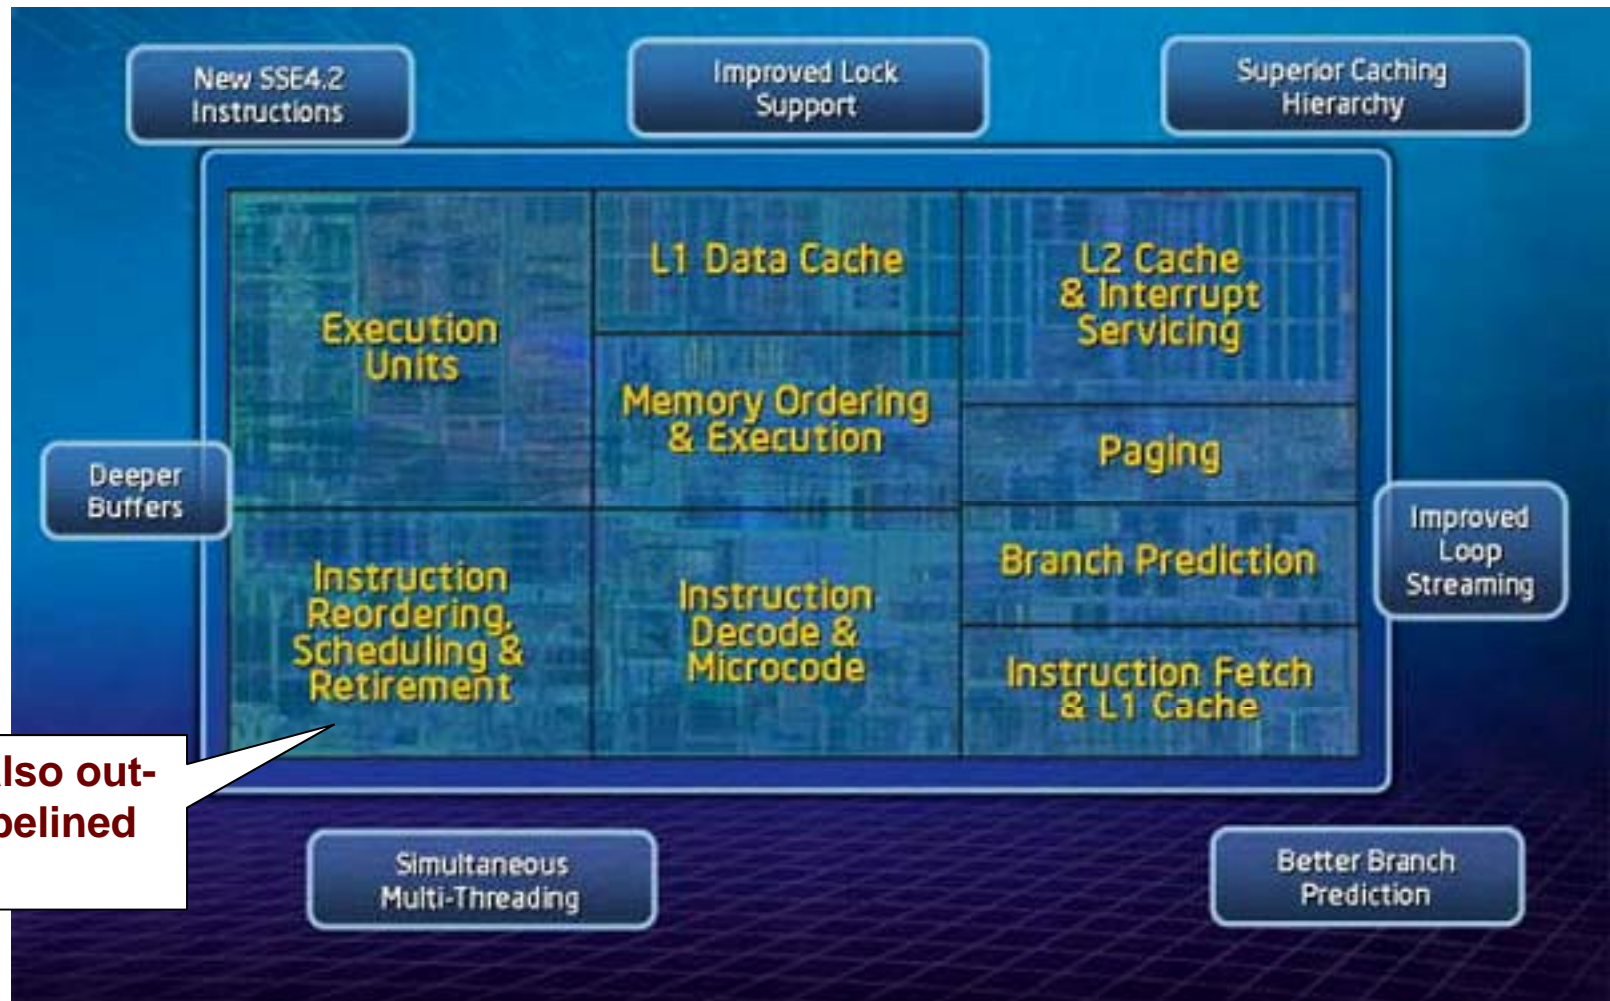
■ Pentium III Microprocessor:



Source: "Digital Design and Computer Architecture" by Harris and Harris, Chapter 7
(read the textbook for other processor architectures and their descriptions)

Visualizing Computer Architectures /2

■ Core i7 Microprocessor Core:



i7 Core is also out-of-order pipelined processor

Source: Intel Core i7 Processor Technical documents

Handling Exceptions /1

■ What are Exceptions?

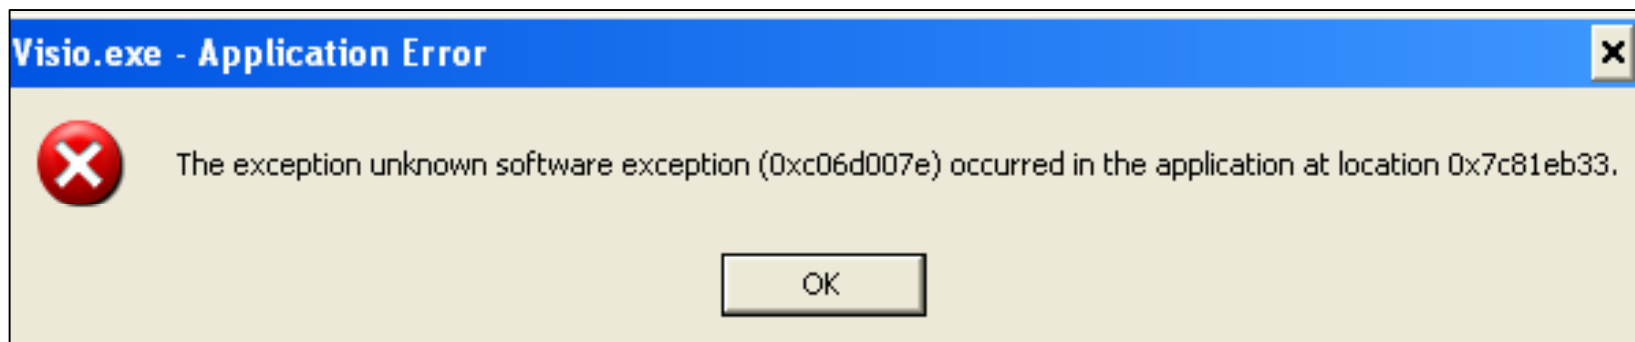
- Unscheduled function call to an exception handler

■ Caused by:

- Hardware (also called interrupt), such as keyboard interrupts
- Software (also called traps), such as undefined instruction

■ When exception occurs, the processor:

- Records cause of exception (Cause register)
- Jumps to exception handler
- Returns to program (EPC register)



Handling Exceptions /2

■ Exception registers are not part of the register file

■ Cause Register

- Records cause of exception
- Coprocessor 0, register 13

■ EPC (Exception PC) Register

- Records PC where exception occurred
- Coprocessor 0, register 14

■ Move from Coprocessor 0

- `mfc0 $t0, Cause`
- Moves contents of `Cause` into `$t0`

mfc0

| | | | | |
|--------|-------|----------|------------|--------------|
| 010000 | 00000 | \$t0 (8) | Cause (13) | 000000000000 |
| 31:26 | 25:21 | 20:16 | 15:11 | 10:0 |

Handling Exceptions /3

■ Exception Causes:

| Exception | Cause |
|------------------------------|------------|
| Hardware Interrupt | 0x00000000 |
| System Call | 0x00000020 |
| Breakpoint / Divide by 0 | 0x00000024 |
| Undefined Instruction | 0x00000028 |
| Arithmetic Overflow | 0x00000030 |

Food for Thought

- **Download and Read Assignment #3 Specifications**
- **Read:**
 - Chapter 6 from the Course Notes
 - Review the material discussed in the lecture notes in more detail
 - Our course schedule follows the material in the Course Notes
 - Recommended: Chapter 7 from the Harris and Harris textbook
- OR
- Recommended: Chapter 4 from the course textbook