

Memory-Mapped Input and Output

Dr. Igor Ivkovic

iivkovic@uwaterloo.ca

[with material from “Computer Organization and Design” by Patterson and Hennessy, and “Digital Design and Computer Architecture” by Harris and Harris, both published by Morgan Kaufmann]

Objectives

- Additional notes on memory management
- Using memory-mapped input and output
- Interconnecting computer components

Practical Virtual Memory Management /1

- **In the examples given before, the size of virtual address space was taken to be 2 GB**
 - In practice, for instance for 32-bit Windows processes, the amount of available virtual memory is 4 GB (or 2^{32} bytes)
 - However, that space is typically divided as 2 GB for user processes, and 2 GB for protected OS processes
 - For 64-bit operating system, the theoretical amount of available virtual memory is 2^{64} , which equals about 16 exabytes
 - In practice, for instance for 64-bit Windows processes, the amount of used virtual memory is much smaller
 - User processes are allocated 8 TB of space, while the OS processes can use up to 248 TB of space

Practical Virtual Memory Management /2

■ **User space vs. system space pages:**

- In Windows, all virtual memory pages in the user space can be paged out (moved) from physical memory to hard disk
- However, not all virtual memory pages in the system space can be paged out from physical memory to hard disk
- **Paged (System) Memory:** Represents the pages that can be paged out from physical memory to hard disk
- **Nonpaged (System) Memory:** Represents the pages that cannot be paged out from physical memory to hard disk

Memory-Mapped Input and Output /1

■ **Accessing Input and Output (I/O) Devices:**

- The processor accesses I/O devices – such as keyboards, monitors, and printers – just like it accesses memory
- **Each I/O device is assigned one or more addresses**
- When that address is detected, data is read from or written to the I/O device instead of memory
- **A portion of the address space is dedicated to I/O devices**
 - Example: addresses in the range 0xFFFF0000 to 0xFFFFFFFF

Memory-Mapped Input and Output /2

■ I/O is mediated by the OS:

- Multiple programs share I/O resources
 - Hence, the OS needs to provide protection and scheduling
- I/O causes asynchronous interrupts
 - Same mechanism as exceptions
- I/O interface needs to be unified
 - OS provides abstractions to programs

■ Device Driver:

- Software that is provided to communicate with an I/O device
- Other programs can call on the code within the device driver to utilize the I/O hardware without knowing its internal details

Memory-Mapped Input and Output /3

- **I/O devices are managed by I/O controller hardware:**
 - Transfers data to/from device
 - Synchronizes operations with software
- **Command registers:**
 - Cause device to do something
- **Status registers:**
 - Indicate what the device is doing and occurrence of errors
- **Data registers:**
 - **Write:** Transfer data to a device
 - **Read:** Transfer data from a device

Memory-Mapped Input and Output /4

■ **I/O Register Mapping:**

- Registers are addressed in the same space as memory
- Address decoder distinguishes between the registers
- OS uses address translation mechanism to make them only accessible to kernel

■ **Address Decoder:**

- Takes an address as input and determines which device/memory communicates with the processor

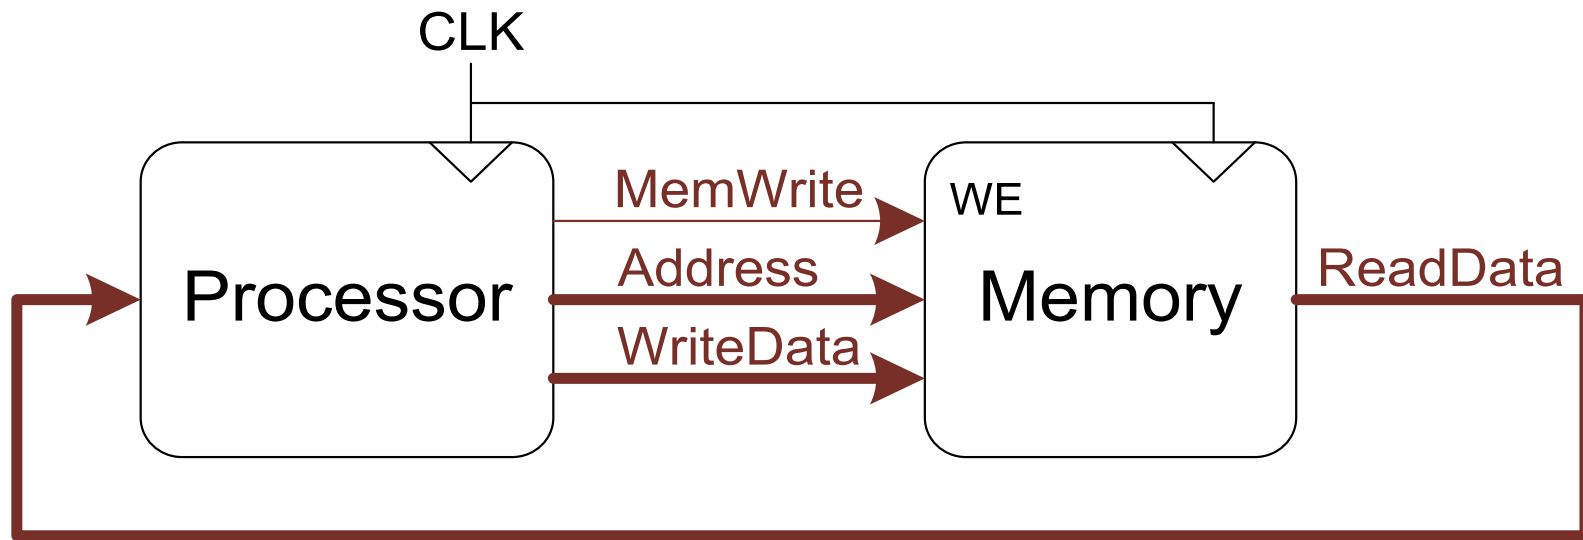
■ **ReadData Multiplexer:**

- A MUX that selects between memory and I/O devices as source of data sent to the processor

Memory-Mapped Input and Output /5

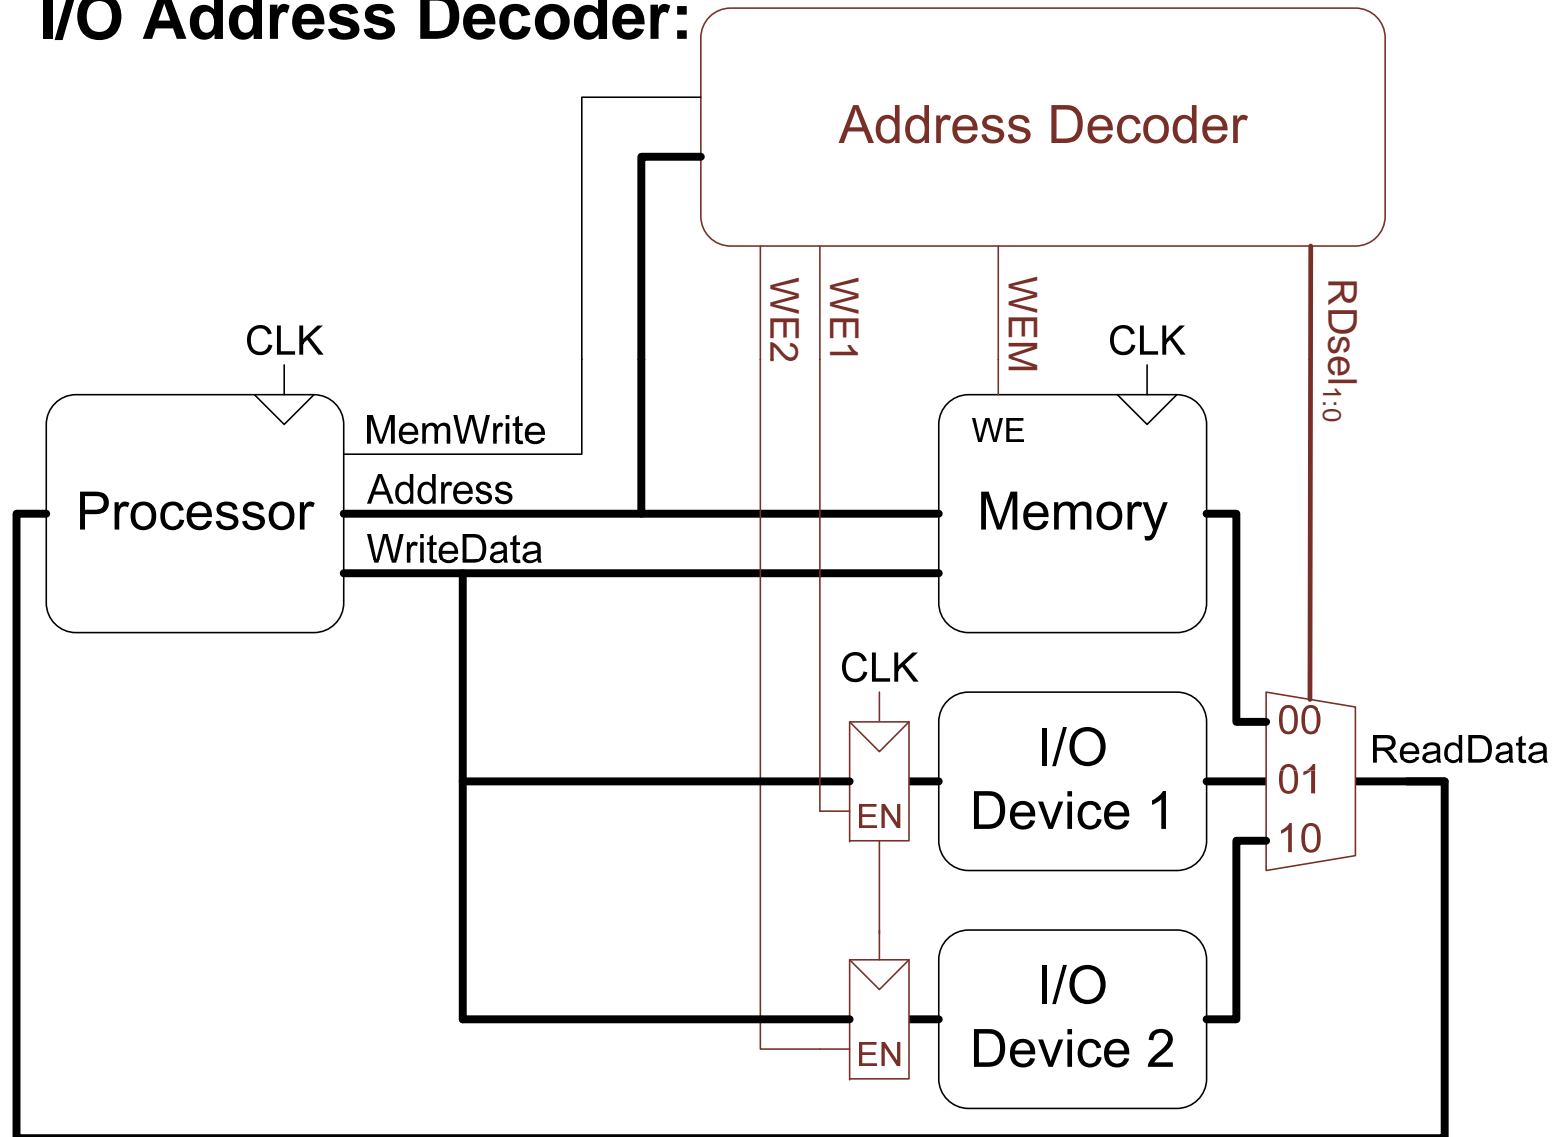
■ Review – CPU-to-Memory Interface:

- What we have covered so far includes an interface between the processor and memory
- If the MemWrite is asserted as 1, Memory writes WriteData to the address specified in the Address input
- If the MemWrite is 0, Memory reads ReadData from the address specified in the Address Input



Memory-Mapped Input and Output /6

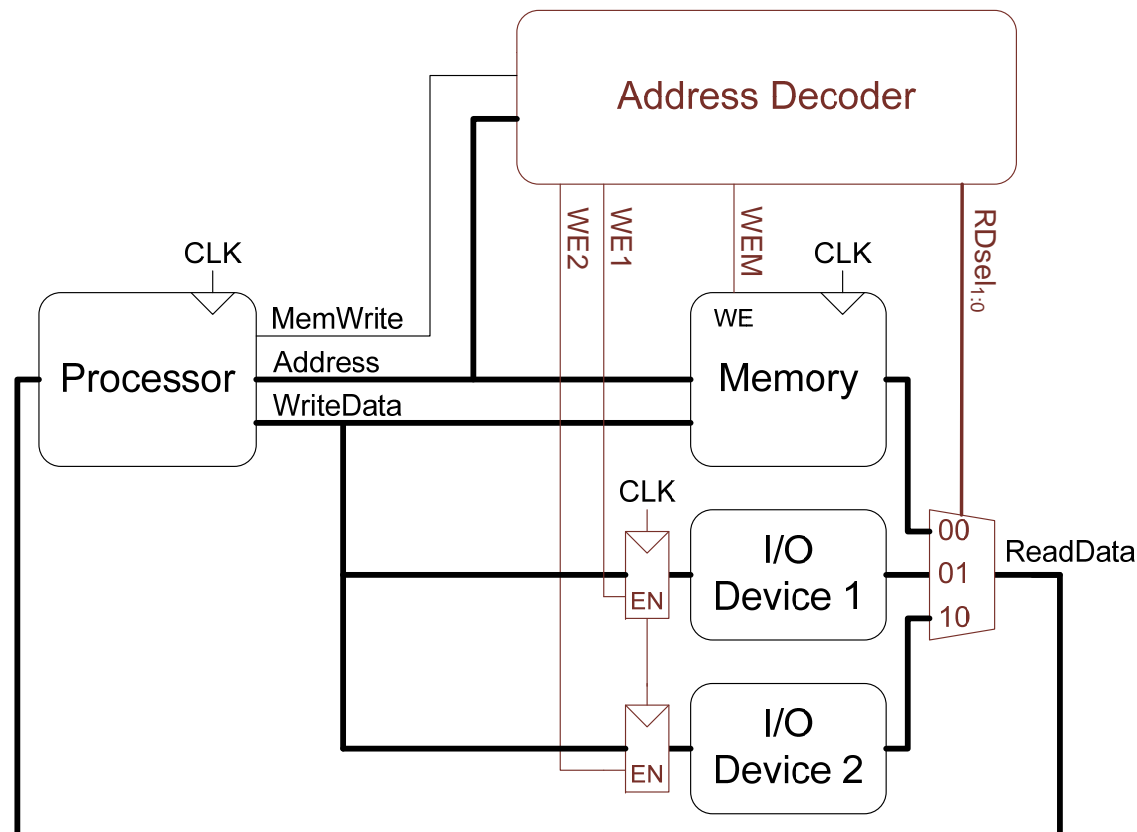
■ I/O Address Decoder:



Memory-Mapped Input and Output /7

■ Address Decoder:

- WE1 and WE2 are Write-Enable signals which enable writes to I/O Device 1 and I/O Device 2, respectively
- RDsel_{1:0} enables read from a specific device or memory



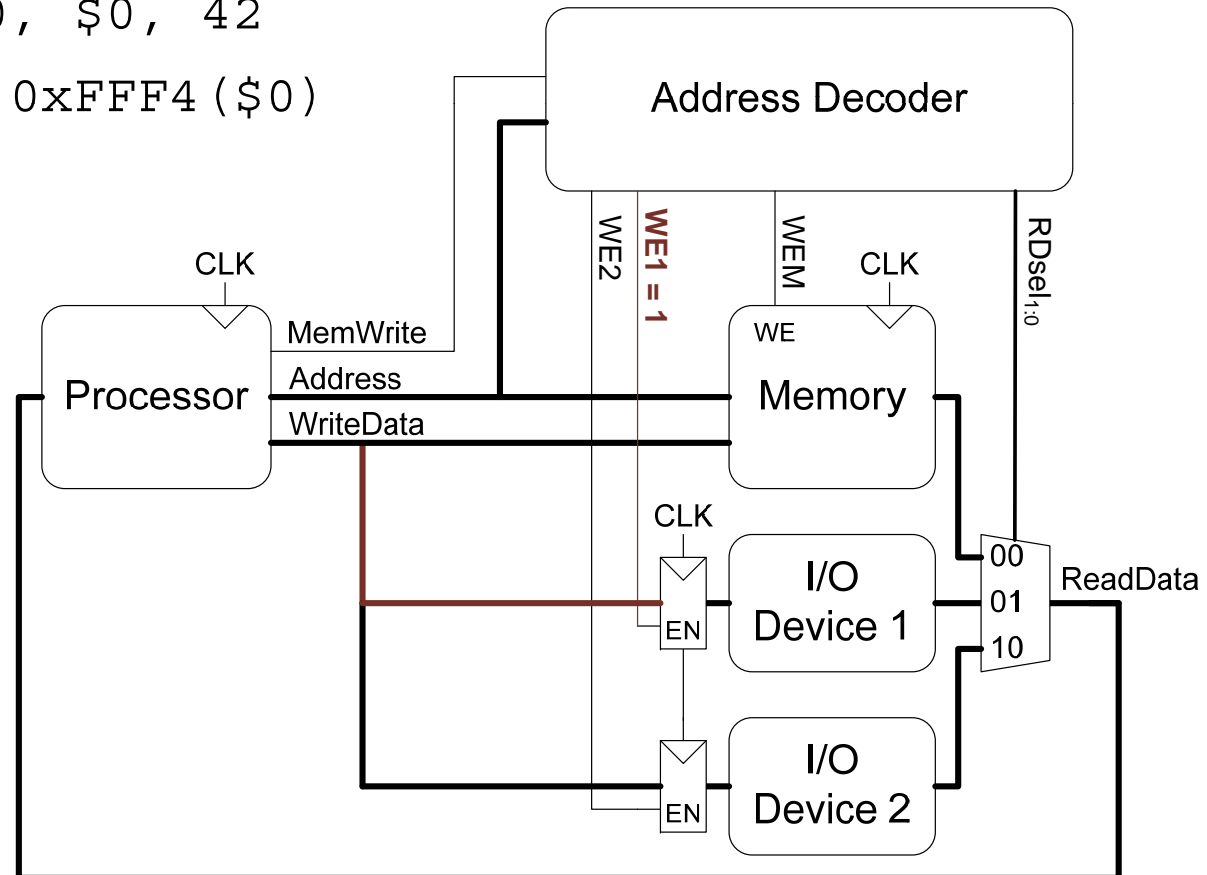
Memory-Mapped Input and Output /8

■ Example:

- Suppose I/O Device 1 is assigned the address 0xFFFFFFFF4
- Write the value 42 to I/O Device 1

```
addi $t0, $0, 42
```

```
sw $t0, 0xFFF4($0)
```

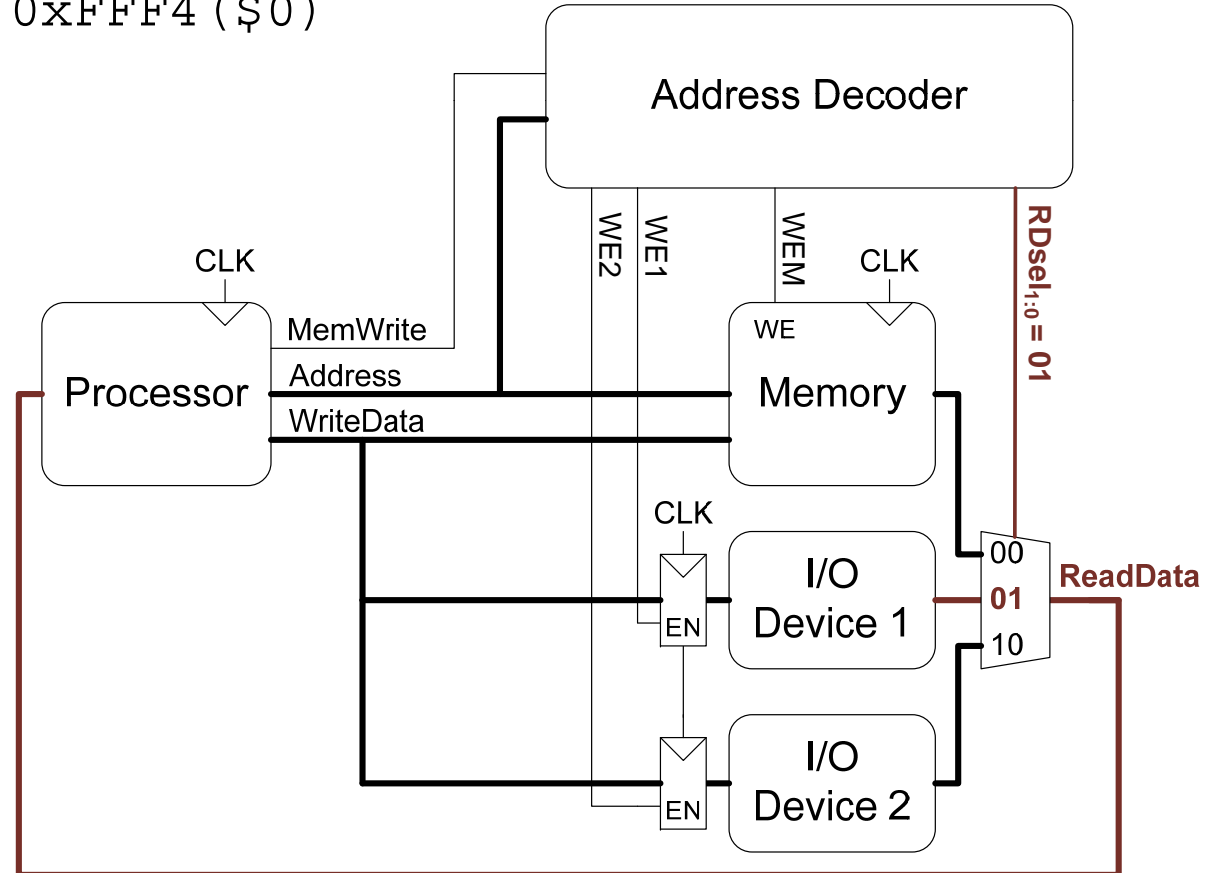


Memory-Mapped Input and Output /9

■ Example:

- Suppose I/O Device 1 is assigned the address 0xFFFFFFF4
- Read the value from I/O Device 1 and place in \$t3

```
lw $t3, 0xFFF4($0)
```



Controlling Memory-Mapped I/O /1

■ I/O Polling:

- Periodically check the I/O status register
- If the device is ready, do the requested operation
- If there is an error, take an appropriate action
- Polling is common in small or low-performance real-time embedded systems, which exhibit
 - Predictable timing, and
 - Low hardware cost
- In other systems, it can waste CPU time

Controlling Memory-Mapped I/O /2

■ **Interrupt-Driven I/O:**

- When a device is ready or error occurs, I/O controller interrupts the CPU
- Interrupt is like an exception (recall Cause register)
 - But interrupts are not synchronized to the instruction execution
 - Can invoke interrupt handler between instructions
 - The cause information often identifies the interrupting device
- **Priority interrupts:**
 - Devices needing more urgent attention get higher priority
 - Can interrupt handler for a lower priority interrupt

Controlling Memory-Mapped I/O /3

- **In polling I/O and interrupt-driven I/O:**
 - CPU transfers data between memory and I/O data registers
 - **This can be time consuming for high-speed devices**
- **Direct Memory Access (DMA):**
 - Instead, OS provides a starting address in memory
 - I/O controller transfers to/from memory autonomously
 - Controller interrupts the CPU on completion or on error

Controlling Memory-Mapped I/O /4

■ **DMA/Cache Interaction:**

- If DMA writes to a main memory block that is cached, the cached copy becomes stale
- If write-back cache has dirty block, and DMA reads the corresponding main memory block, DMA reads stale data

■ **Hence, we need to ensure cache coherence**

- Flush blocks from cache if they will be used for DMA
- Or use non-cacheable memory locations for I/O

Controlling Memory-Mapped I/O /5

■ **DMA/VM Interaction:**

- OS uses virtual addresses for memory, but DMA blocks may not be contiguous in physical memory
- Should the DMA use virtual addresses?
 - This would require I/O controller to do the VM translation

■ **Instead, if DMA uses physical addresses**

- Need to break transfers into page-sized data fragments
- Or allocate contiguous physical pages for DMA

Controlling Memory-Mapped I/O /6

■ **Measuring I/O Performance:**

- I/O performance depends on
 - **Hardware:** CPU, memory, controllers, buses
 - **Software:** Operating system, application software
 - **Workload:** Request rates and patterns
- **I/O system design can trade response time for throughput**
 - Measurements of throughput often done with constrained response-time

Interconnecting Computer Components /1

- **CPU hardware requires interconnections between**

- CPU, memory, I/O controllers

- **(System) Bus:**

- A shared communication channel
- Parallel set of wires for data and synchronization of data transfer
- Can become a performance bottleneck
- Performance limited by physical factors, such as wire length, and number of connections

Interconnecting Computer Components /2

- **Data lines:**

- Carry address and data
- Multiplexed or separate

- **Control lines:**

- Indicate data type, synchronize transactions

- **Synchronous line:**

- Uses a bus clock

- **Asynchronous line:**

- Uses request/acknowledge control lines for handshaking

Interconnecting Computer Components /3

■ **Processor-Memory buses:**

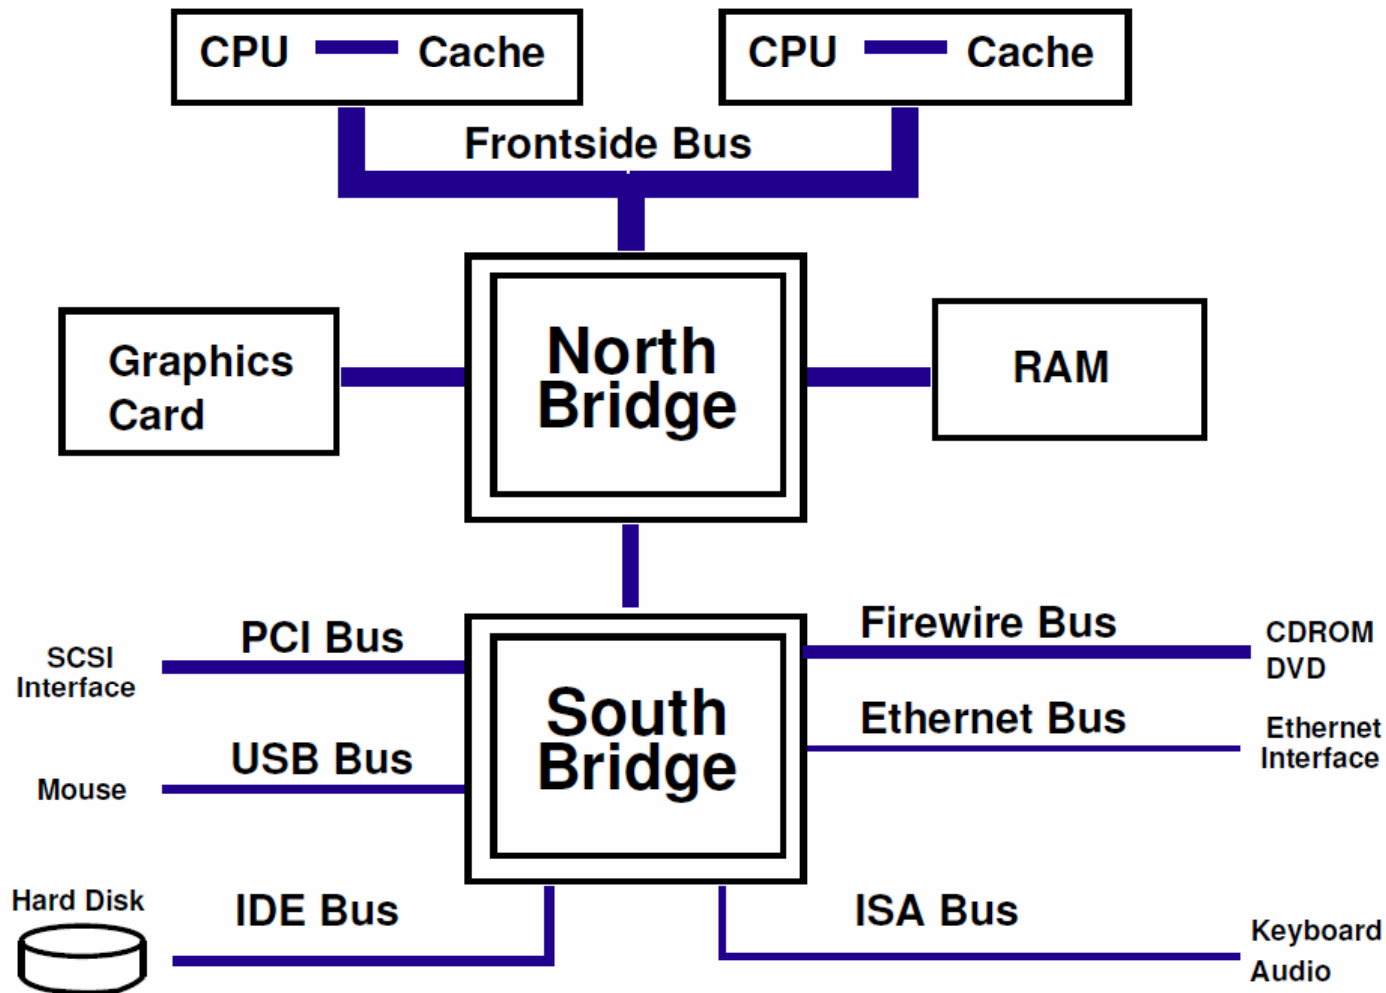
- Short in length and high speed in performance
- Design of the memory bus is setup in a manner that matches the memory organization (e.g., presence of L1/L2/L3 caches)

■ **I/O buses:**

- Longer in length and allow multiple connections
- Specified by standards for interoperability
- **Connect to processor-memory bus through a bridge**
- **North Bridge (Memory Controller):** In charge of controlling the transfers between the processor and main memory (RAM)
- **South Bridge (I/O Controller or I/O Controller Hub (ICH)):** Controls the transfers between peripheral devices

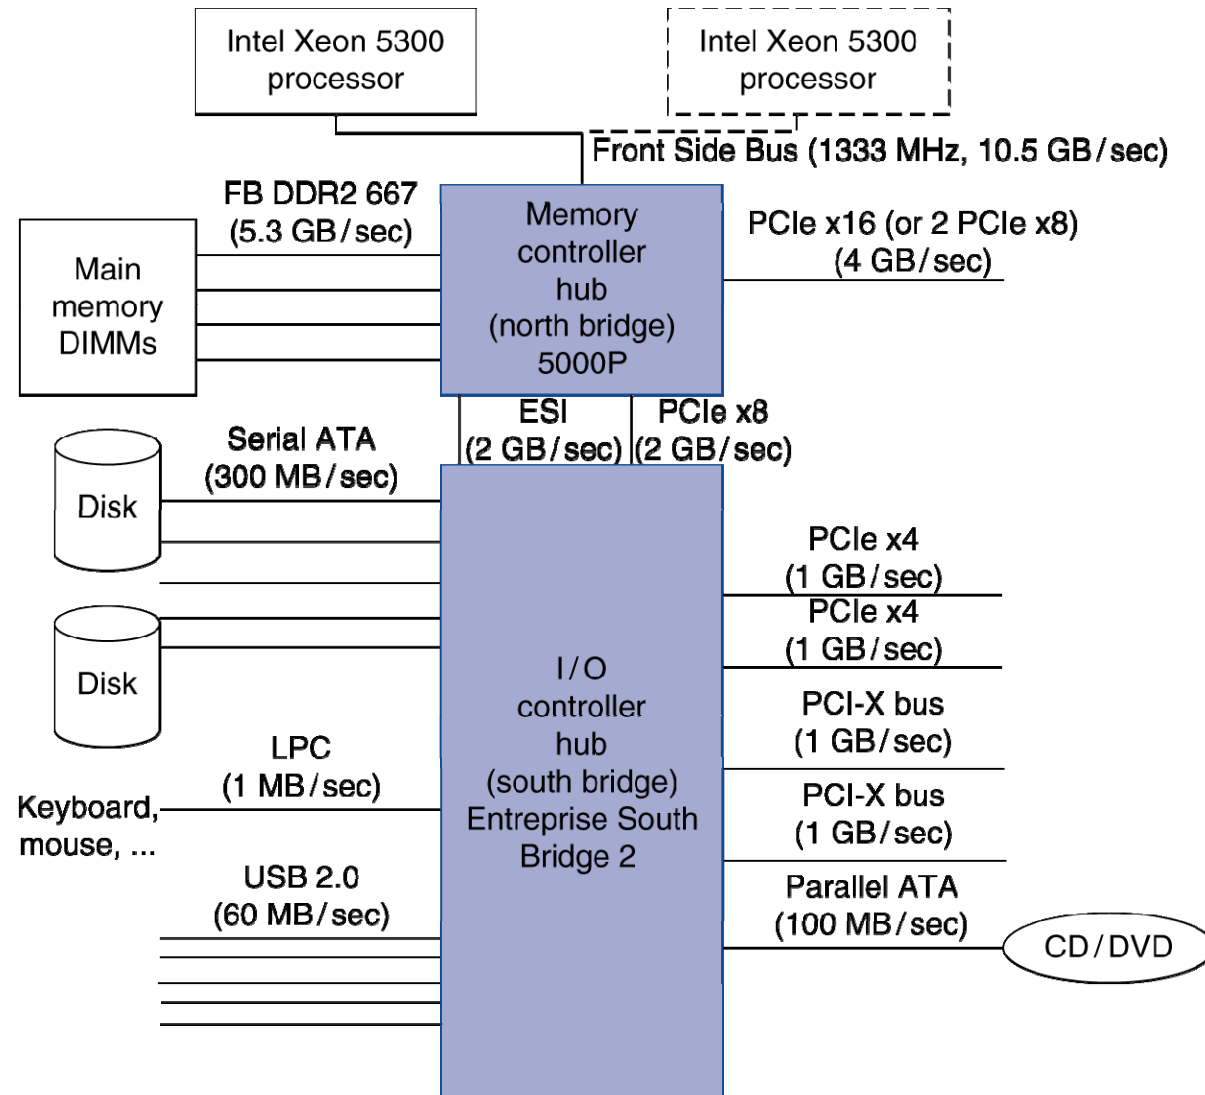
Interconnecting Computer Components /4

■ Example – Pentium Bus



Interconnecting Computer Components /5

■ Example – Typical x86 PC I/O System:



Interconnecting Computer Components /6

■ I/O Bus Examples:

	Firewire	USB 2.0	PCI Express	Serial ATA	Serial Attached SCSI
Intended use	External	External	Internal	Internal	External
Devices per channel	63	127	1	1	4
Peak bandwidth	50MB/s or 100MB/s	0.2MB/s, 1.5MB/s, or 60MB/s	250MB/s/lane 1x, 2x, 4x, 8x, 16x, 32x	300MB/s	300MB/s
Hot pluggable	Yes	Yes	Depends	Yes	Yes
Max length	4.5m	5m	0.5m	1m	8m
Standard	IEEE 1394	USB Implementers Forum	PCI-SIG	SATA-IO	INCITS TC T10

Food for Thought

- **Download and Read Assignment #5 Specifications**
- **Read:**
 - Chapter 8 from the Course Notes
 - Review the material discussed in the lecture notes in more detail
 - Our course schedule follows the material in the Course Notes
 - Recommended: Chapter 8 from the Harris and Harris textbook
- OR
- Recommended: Chapter 6 from the course textbook